

How to Avoid Making a Billion-Dollar Mistake: Type-Safe Data Plane Programming with SafeP4

Anonymous Authors

Abstract

The P4 programming language offers high-level, declarative abstractions that bring the flexibility of software to the domain of networking. Unfortunately, the main abstraction used to represent packet data in P4—header types—lacks basic safety guarantees. Over the last few years, experience with an increasing number of P4 programs has shown the risks of the unsafe approach, which often leads to subtle software bugs.

This paper proposes SAFE P4, a domain-specific language for programmable data planes in which all packet data is guaranteed to have a well-defined meaning and satisfy essential safety guarantees. We equip SAFE P4 with a formal semantics and a static type system that statically guarantees header validity—a major source of safety bugs according to our analysis of real-world P4 programs. Statically ensuring header validity is challenging because the set of valid headers can be modified at runtime, making it a dynamic program property. Our type system achieves static safety by using a form of path-sensitive reasoning that tracks dynamic information from conditional statements, routing tables, and the control plane. Our empirical evaluation shows that SAFE P4’s type system can effectively eliminate common failures in many real-world programs.

2012 ACM Subject Classification Software and its engineering → Formal language definitions, Networks → Programming interfaces

Keywords and phrases P4, data plane programming, type systems

Digital Object Identifier 10.4230/LIPIcs.CVIT.2016.23

1 Introduction

I couldn’t resist the temptation to put in a null reference [...] This has led to innumerable errors, vulnerabilities, and system crashes, which have probably caused a billion dollars of pain and damage in the last forty years.

—Tony Hoare

Modern languages offer high-level abstractions that provide numerous benefits to programmers. Features such as type systems, structured control flow, objects, modules, etc. make it possible to express rich computations in terms of high-level structures and algorithms rather than machine-level code. Increasingly, many languages also offer fundamental safety guarantees—e.g., well-typed programs do not go wrong [21]—that make entire categories of programming errors simply impossible.

Unfortunately, although computer networks are critical infrastructure, providing the essential communication fabric that underpins nearly all modern systems, most network devices today are programmed using low-level languages that lack even basic safety guarantees. Unsurprisingly, networks are unreliable and remarkably insecure—e.g., the first step in the majority of cyberattacks involves compromising a router or other device [24, 17].

Over the past decade, there has been a remarkable shift to more flexible platforms in which the functionality of the network is specified in software. Early efforts related to



© ; licensed under Creative Commons License CC-BY

42nd Conference on Very Important Topics (CVIT 2016).

Editors: John Q. Open and Joan R. Access; Article No. 23; pp. 23:1–23:45



Leibniz International Proceedings in Informatics

LIPIcs Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

software-defined networking (SDN) [19, 6], focused on the control software that computes routes, balances load, and enforces security policies, and modeled the data plane as a simple pipeline of routing tables operating on a fixed set of packet formats. However, there has been recent interest in allowing the functionality of the data plane itself to be specified as a program—e.g., to implement new protocols, make more efficient use of hardware resources, or even relocate application-level functionality into the network [13, 12]. In particular, the P4 language [4] enables the functionality of a data plane to be programmed in terms of declarative abstractions such as header types, packet parsers, routing tables, and structured control flow that a compiler maps down to an underlying target device.

Unfortunately, while a number of P4’s features were clearly inspired by designs found in modern languages, the central abstraction for representing packet data, header types, lacks basic safety guarantees. To a first approximation, a P4 header type can be thought of as a record with a field for each component of the header. For example, the header type for an IPv4 packet, would have a 4-bit version field, an 8-bit time-to-live field, two 32-bit fields for the source and destination addresses, and so on.

Unfortunately, according to the P4 language specification, instances of a header type may either be valid or invalid, and reading or writing an invalid header yields an undefined result. In practice, reading or writing an invalid header can lead to a variety of problems including dropping the packet when it should be forwarded, or even leaking information from one packet to the next! It also breaks portability, since the same program may behave differently when executed on different targets.

The choice to model the semantics of header types in an unsafe way was intended to make the language easier to implement on high-speed routers, which typically have limited amounts of memory. A typical P4 program might specify behavior for several dozen different protocols, but any particular packet is likely to contain only a small handful of headers. It follows that if the compiler only needs to represent the valid headers at run-time, then memory requirements can be reduced. However, while it may have benefits for language implementers, the design is a disaster for programmers—it repeats Hoare’s “billion-dollar mistake,” and bakes an unsafe feature deep into the design of a language that has the potential to become the de-facto standard in a (multi) billion-dollar industry.

This paper investigates the design of a domain-specific language for programmable data planes in which all packet data is guaranteed to have a well-defined meaning and satisfy basic safety guarantees. In particular, we present **SAFE**P4, a language with a precise semantics and a static type system that guarantees the validity of all headers read or written by the program. Although the type system is based on standard features, there are several aspects of its design that stand out. First, to facilitate tracking dependencies between headers—e.g. if the TCP header is valid, then the IPv4 will also be valid—SAFE P4 has an expressive algebra of types that keeps tracks of fine-grained validity information. Second, to enable SAFE P4 to accommodate the growing collection of extant P4 programs with only modest modifications, it uses a path-sensitive type system that incorporates information from conditional statements, routing tables, and the control plane to precisely track validity.

To evaluate our design for SAFE P4, we formalized the language and its type system in a core calculus and proved the usual progress and preservation theorems. We also implemented the SAFE P4 type system in an OCaml prototype, **P4CHECK**, and applied it to a suite of open-source programs found on GitHub such as **switch.p4**, a large P4 program that implements the features found in modern data center switches (specifically, it includes over four dozen different switching, routing, and tunneling protocols, as well as multicast, access control lists, among other features). We categorize common failures and, for programs that

fail to type check, identify the root causes and apply repairs to make them well typed. We find that most programs can be repaired with low effort from programmers, typically by applying a modest number of simple repairs.

Overall, the main contributions of this paper are as follows:

- We propose SAFEP4, a type-safe enhancement of the P4 language that eliminates all errors related to header validity.
- We formalize the syntax and semantics of SAFEP4 in a core calculus and prove that the type system is sound.
- We implement our type checker in an OCaml prototype, P4CHECK.
- We evaluate our type system empirically on over a dozen real-world P4 programs and identify common errors and repairs.

The rest of this paper is organized as follows. Section 2 provides a more detailed introduction to P4 and elaborates on the problems this work addresses. Section 3 presents the design, operational semantics and type system of SAFEP4 and reports our type safety result. The results of evaluating SAFEP4 in the wild are presented in Section 4. Section 5 surveys related work and Section 6 summarizes the paper and outlines topics for future work.

2 Background and Problem Statement

This section introduces the main features of P4 and highlights the problems caused by the unsafe semantics for header types using examples.

2.1 P4 Language

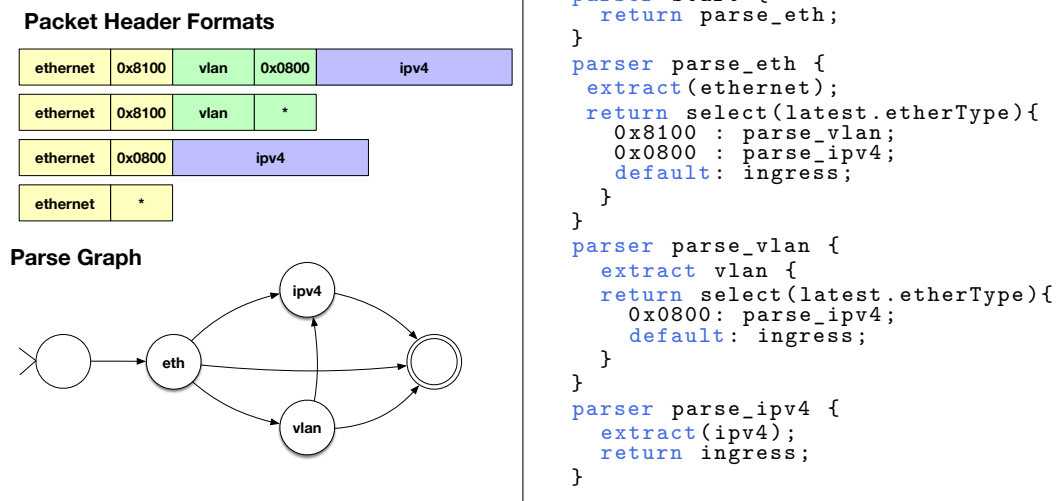
A P4 program comprises header type definitions, a parse graph, table and action declarations, and control flow specifying the order in which tables are applied to network packets.

Header Types and Instances The internal structure of each packet header are specified using P4's header types. For example, the first few lines of the following snippet of code:

```
header_type ethernet_t {
  fields {
    dstAddr: 48;
    srcAddr: 48;
    etherType: 16;
  }
}
header ethernet_t ethernet;
header ethernet_t inner_ethernet;
```

declare a header type (`ethernet_t`) for the Ethernet protocol with fields `dstAddr`, `srcAddr`, and `etherType`. The integer literals indicate the bit width of each field. The next two lines declare two `ethernet_t` instances, each with a distinct name (`ethernet` and `inner_ethernet`) and global scope. Ordinary packets usually have a single Ethernet header, but a tunneling protocol might maintain a second header for the nested packet.

Parsers The order in which packet headers are read from the stream of bits that represents the input packet is specified using P4's parsers, which are modeled using finite state machines. The code within each state may extract bits from the input stream, modify header instances, conditionally branch, and transition either to another state or into the `ingress` control, which has the effect of accepting the packet. Figure 1 depicts a visual representation of a parse graph for three common headers: Ethernet, VLAN, and IPv4. The instance `ethernet`



■ **Figure 1** (Left) Header formats and parse graph that extracts an Ethernet header optionally followed by VLAN and/or IPv4 headers. (Right) P4 code implementing the same parser.

```

table forward {
    reads {
        ipv4 : valid;
        vlan : valid;
        ipv4.dstAddr: ternary;
    }
    actions = {
        nop;
        next_hop;
        remove;
    }
    default_action : nop();
}
    
```

Runtime Contents of forward

Pattern			Action	
ipv4	vlan	ipv4.dstAddr	Name	Data
1	0	10.0.0.0/24	next_hop	<i>m</i>
0	1	0.0.0.0/0	remove	

■ **Figure 2** P4 tables.

131 is extracted first, optionally followed by a `vlan` instance, or an `ipv4` instance, or both.
 132 Extracting into an header instance populates its fields with bits from the input stream,
 133 advances the stream, and then marks the instance as valid.

134 **Tables and Actions** In most P4 programs, the bulk of the processing of each packet is
 135 performed using tables and actions. A table is defined in terms of (i) state it reads to
 136 determine a matching entry (if any), (ii) actions it may execute, and (iii) an optional default
 137 action it executes if no matching entry is found. For example, Figure 2 declares the table
 138 `forward` with a `reads` declaration that checks whether the `ipv4` and `vlan` instances are
 139 valid, and performs ternary matching on the `dstAddr` field of the `ipv4` instance. Tables are
 140 populated at run-time by the control plane, with entries that contain a pattern, an action,
 141 and action data. The pattern specifies the bits that should be compared against the values
 142 read in the table, the action is the name of a function, defined elsewhere in the P4 program,
 143 and the action data are values that serve as the arguments to that function.

144 Operationally, to process a packet, a table firsts scans its entries to locate a matching
 145 entry. If it finds an entry, the table is said to hit, and it executes the associated action.

```

146 action next_hop(src, dst) {
147     modify_field(ethernet.srcAddr, src);
148     modify_field(ethernet.dstAddr, dst);
149     subtract_from_field(ipv4.ttl, 1);
150 }
151
152 action remove() {
153     modify_field(ethernet.etherType,
154                 vlan.etherType);
155     remove_header(vlan);
156 }

```

■ **Figure 3** P4 actions.

146 Otherwise, if there are no matching entries, the packet is said to miss in the table and it
 147 executes the `default_action` (or an implicit no-op if the table lacks a default). For example,
 148 in Figure 2, the `forward` table is shown populated with two rules. The first rule tests whether
 149 `ipv4` is valid, `vlan` is invalid and then applies `next_hop` to an argument m (which stands
 150 for an Ethernet MAC address). The second rule checks that `ipv4` is invalid, then that `vlan`
 151 is valid, and then performs a ternary match using the all-wildcard pattern `0.0.0.0/0`.

152 Actions are functions containing sequences of primitive commands that perform operations
 153 such as adding and removing headers, assigning a value to a field, adding one field to another,
 154 etc. For example, Figure 3 depicts the `next_hop` and `remove` actions. The `next_hop` action
 155 assigns its argument `src`, which is provided by the control plane, to the `srcAddr` field of
 156 the `ethernet` instance while `remove` copies the `etherType` field from the `vlan` instance to
 157 the `ethernet` instance, and then removes the `vlan` instance. The `modify_field($h.f$, e)`
 158 primitive evaluates e and stores the result in header field $h.f$ —e.g., `next_hop` decrements
 159 `ipv4.ttl` by one and saves the result. The `remove_header` primitive invalidates the header
 160 provided as an argument, and behaves like a no-op if the header is already invalid.

161 **Control** Control blocks use standard control-flow constructs to execute a pipeline of tables
 162 in sequence. They manage the order and conditions under which each table is executed. The
 163 `apply` command executes a table and conditionals branch on a boolean expression such as
 164 the validity of a header instance.

```

165 control ingress {
166     if(valid(ipv4) or valid(vlan)) {
167         apply(forward);
168     }
169 }

```

170 This code applies the `forward` table only if one of `ipv4` or `vlan` is valid.

171 2.2 Common Bugs in P4 Programs

172 Having introduced the basic features of P4, we now present five categories of bugs found
 173 in open-source programs that arise due to reading and writing invalid headers—the main
 174 problem that SAFE P4 addresses. There is one category for each of the following syntactic
 175 constructs: (1) parsers, (2) controls, (3) table reads, (4) table actions, and (5) default actions.

176 To identify the bugs we surveyed a benchmark suite of 15 research and industrial P4
 177 programs that are publicly available on GitHub and compile to the BMv2 [22] backend.
 178 Later, in Section 4, we will report the number of occurrences of each of these categories in
 179 our benchmark suite detected by our approach.¹

¹ We focus on P4₁₄ programs in this paper, but the issues we address also persist in the latest version of the language, P4₁₆. We did not consider P4₁₆ due to the smaller number of programs that are available, and for the pragmatic reason that our tool reuses an existing P4₁₄ front-end.

<pre> /* UNSAFE */ parser parse_ethernet { extract(ethernet); return select(latest.etherType) { 0x0800 : parse_ipv4; default : ingress; } } parser parse_ipv4 { extract(ipv4); return select(latest.protocol) { 6 : parse_tcp; default : ingress; } } parser parse_tcp { extract(tcp); return ingress; } </pre>	<pre> /* SAFE */ parser_exception unsupported { parser_drop; } parser parse_ethernet { extract(ethernet); return select(latest.etherType) { 0x0800 : parse_ipv4; default : parser_error unsupported; } } parser parse_ipv4 { extract(ipv4); return select(latest.protocol) { 6 : parse_tcp; default : parser_error unsupported; } } control ingress { if (tcp.syn == 1 and tcp.ack == 1) { ... } } </pre>
--	--

■ **Figure 4** Left: unsafe code in NETHCF; Right: our type-safe fix; Bottom: common code.

2.2.1 Parser Bugs

The first class of errors is due to the parser being too conservative about dropping malformed packets, which increases the set of headers that may be invalid in the control pipeline. In most programs, the parser chooses which headers to **extract** based on the fields of previously-extracted headers. For example, after extracting the IPv4 header, it is common to branch on the **protocol** field to decide whether to extract the TCP (0x06) or UDP (0x11) header. However, these alternatives do not capture all possible values for the **protocol** field—e.g., consider an OSPF packet (0x59). A common programming pattern is to add a **default** case to ensure that all possible packets are handled. However, the oft-forgotten next step is to drop the offending packet in the **ingress** pipeline! Failure to do so leads to undefined behavior, whenever the programmer attempts to access the TCP or UDP header.

An example from the NETHCF [31, 2] codebase illustrates this bug. NETHCF is a research tool designed to combat TCP spoofing. The program detects and drops spoofed packets by learning allowable hop-count values for the IPv4 address of each host and discarding packets from those hosts that have invalid hop-counts. As shown in Figure 4, the parser handles IPv4-TCP packets and redirects all other packets to the **ingress** control. Unfortunately, the **ingress** control executes no check on whether **tcp** is valid before accessing **tcp.syn** to check whether it is equal to 1. This is unsafe since **tcp** is not guaranteed to be valid.

To fix this bug, we can define a parser exception, **unsupported**, with an handler that drops packets, thereby protecting the **ingress** from having to handle unexpected packets.

2.2.2 Control Bugs

Another common bug occurs when a table is executed in a context in which the instances referenced by that table are not guaranteed to be valid. This bug can be seen in the open-source code for NETCACHE [11, 13], a P4 system that implements a load-balancing cache.

<pre> 204 /* UNSAFE */ 205 control ingress { 206 207 process_cache(); 208 process_value(); 209 210 apply(ipv4_route); 211 } </pre>	<pre> 204 /* SAFE */ 205 control ingress { 206 if (valid(nc_hdr)) { 207 process_cache(); 208 process_value(); 209 } 210 apply(ipv4_route); 211 } </pre>
<pre> 212 control process_cache { 213 apply(check_cache_exist); 214 ... 215 } </pre>	<pre> 212 table check_cache_exist { 213 reads { nc_hdr.key } 214 actions { ... } 215 } </pre>

■ **Figure 5** Left: unsafe code in NETCACHE; Right: our type-safe fix; Bottom: Common code

204 The parser for NETCACHE reserves a specific port (8888) to handle its special-purpose traffic,
 205 a condition that is built into the parser, which extracts `nc_hdr` (i.e., the `NETCACHE`-specific
 206 header) only when UDP traffic arrives from port 8888. Otherwise it performs standard L2 and
 207 L3 routing. Unfortunately, the `ingress` control node tries to access `nc_hdr` before checking
 208 that it is valid! Specifically, the `reads` declaration for `check_cache_exists` table, which is
 209 executed first, presupposes that `nc_hdr` is valid. The invocation of the `process_value` table
 210 (not shown) contains another instance of the same bug.

211 To fix these bugs, we can wrap the calls to `process_cache` and `process_value` in an
 212 conditional that checks the validity of the header `nc_hdr`. This ensures that `nc_hdr` is valid
 213 when `process_cache` refers it.

214 2.2.3 Table Reads Bugs

215 A similar bug arises in programs that contain tables that first match on the validity of
 216 certain header instances before matching on the fields of those instances. The advantage
 217 of this approach is that multiple types of packets can be processed in a single table, which
 218 saves memory. However, if implemented incorrectly, it can lead to a bug, in which the `reads`
 219 declaration matches on the `exact` bits in a field from a header that may not be valid!

220 The `switch.p4` program is a “realistic production switch” [16] developed by Barefoot
 221 Networks, meant to be used “as-is, or as a starting point for more advanced switches” [16].
 222 It provides many core network features, such as L2 switching, L3 routing, LAG, ECMP,
 223 VLAN, NVGRE, Geneve, GRE, ACL, and MPLS, among others [16, 23]. An archetypal
 224 example of table reads bugs is the `port_vlan_mapping` table of SWITCH (Figure 6). This
 225 table is invoked in a context, where it is not known which of the VLAN tags is valid, despite
 226 containing references to both `vlan_tag_0` and `vlan_tag_1` in its `reads` declaration, so
 227 the programmer has guarded the references to `vlan_tag_0` with keys that test the
 228 validity of `vlan_tag_0`, for $i = 1, 2$.

229 Guarding the accesses with valid matches in this way allows the control plane to install
 230 rules with keys such as `(0, 0, 1, 2)`, which checks the validity of each header before accessing
 231 them to perform the match. However, it is impossible, with the table as it is written, for the
 232 control plane to install a rule that will avoid reading the value of an invalid header. The first
 233 match that is performed at runtime is to verify that the `vlan_tag_0` instance is invalid.
 234 Then the `exact` match kind will try and check that `vlan_tag_0.vid = 0`, even though the
 235 instance is invalid! This attempt to access an invalid header results in undefined behavior,
 236 and is therefore a bug.

<pre> 237 /* UNSAFE */ 238 table port_vlan_mapping { 239 reads { 240 vlan_tag_[0] : valid; 241 vlan_tag_[0].vid : exact; 242 vlan_tag_[1] : valid; 243 vlan_tag_[1].vid : exact; 244 } ... 245 } </pre>	<pre> 237 /* SAFE */ 238 table port_vlan_mapping { 239 reads { 240 vlan_tag_[0] : valid; 241 vlan_tag_[0].vid : ternary; 242 vlan_tag_[1] : valid; 243 vlan_tag_[1].vid : ternary; 244 } ... 245 } </pre>
--	--

■ **Figure 6** Left: a table in `switch.p4` with unprotected conditional reads; Right: our type-safe fix.

237 Its worthy to note that this code is not actually buggy on some targets—in particular,
 238 if invalid headers are initialized with 0. However, this doesn’t conform to the language
 239 specification, and therefore isn’t portable to other hardware devices.

240 The naive solution to fix this bug is to refactor the table into four different tables (one
 241 for each combination of validity bits) and then check the validity of each header before the
 242 tables are invoked. This is a perfectly safe fix, but it can result in an exponential blowup
 243 in the number of tables, which is clearly undesirable for efficiency reasons, and because it
 244 complicates the control plane.

245 Fortunately, rather than factoring the table into four tables, we can replace the `exact`
 246 match-kinds with `ternary` match-kinds, which permit matching with wildcards. Since
 247 matching a wildcard does not evaluate the expression, then match rules like `(0, *, 0, *)` will
 248 cause the switch to safely skip evaluation of `vlan_tag_[i].vid`.

249 In order for this solution to typecheck, we need to assume that the control plane is
 250 well-behaved—i.e. that it will install wildcards for the `ternary` matches whenever the
 251 header is invalid. In our implementation, we print a warning whenever we make this kind of
 252 assumption so that the programmer can confirm that the control plane is well-behaved.

253 2.2.4 Table Action Bugs

254 Another prevalent bug, in our experience, arises when distinct actions in a table require
 255 different (and possible mutually exclusive) headers to be valid. This can lead to two problems:
 256 (i) the control plane can populate the table with unsafe match-action rules, and (i) there may
 257 be no validity checks that we can add to the control to make all of the actions typecheck.

258 The `fabric_ingress_dst_lkp` table (Figure 7) in `switch.p4` provides an example of this
 259 misbehavior. The `fabric_ingress_dst_lkp` table reads the value of `fabric_hdr.dstDevice`
 260 and then invokes one of several actions: `term_cpu_packet`, `term_fabric_unicast_packet`,
 261 or `term_fabric_multicast_packet`. Respectively, these actions require the `fabric_hdr_cpu`,
 262 `fabric_hdr_unicast`, and `fabric_hdr_multicast` (respectively) headers to be valid. Un-
 263 fortunately the validity of these headers is mutually exclusive².

264 Since `fabric_header_cpu`, `fabric_header_unicast`, and `fabric_header_multicast`
 265 are mutually exclusive, there is no context that makes this table safe. The only facility the ta-
 266 ble provides to determine which action should be called is `fabric_hdr.dstDevice`. However,
 267 the parser chooses which header to parse based on the value of `fabric_hdr.packetType`, and
 268 the portion of the program that precedes the call to `fabric_ingress_dst_lkp` does not estab-
 269 lish any relationship between `fabric_hdr.dstDevice` and `fabric_hdr.dst_device`. In fact

² There are other actions in the real `fabric_ingress_dst_lkp`, but these three actions demonstrate the core of the problem.

<pre> 270 /* UNSAFE */ 271 table fabric_ingress_dst_lkp { 272 reads { 273 fabric_hdr.dstDevice : exact; 274 } 275 276 actions { 277 term_cpu_packet; 278 term_fabric_unicast_packet; 279 term_fabric_multicast_packet; 280 } 281 } </pre>	<pre> 270 /* SAFE */ 271 table fabric_ingress_dst_lkp { 272 reads { 273 fabric_hdr.dstDevice : exact; 274 fabric_hdr_cpu : valid; 275 fabric_hdr_unicast : valid; 276 fabric_hdr_multicast : valid; 277 } 278 actions { 279 term_cpu_packet; 280 term_fabric_unicast_packet; 281 term_fabric_multicast_packet; 282 } 283 } </pre>
---	---

■ **Figure 7** Left: unsafe code in `switch.p4`; Right: our type-safe fix.

270 the program does not even reference these locations before the call to `fabric_ingress_dst_lkp`.
 271 Hence, the correctness of this table relies on well-formed input packets, which is not consistent
 272 with real switches, which can receive any sequence of bits that arrive “on the wire.”

273 To fix this bug there are two possible solutions: (1) refactor the table into three tables
 274 whose applications are guarded with validity checks for the required headers, or (2) include
 275 validity matches in the reads declaration. In general, to avoid the exponential blowup in
 276 the number of on-switch tables, we proceed with option (2) as shown on the right side of
 277 Figure 7.

278 In order to type check this solution, we need to make an assumption about the way the
 279 control plane will populate the table. Concretely, if an action a only typechecks if a header h
 280 is valid, and h is not necessarily valid when the table is applied, we assume that the control
 281 plane will only call a if h is matched as valid. For example, `fabric_hdr_cpu` is not known to
 282 be valid when (the fixed version of) `fabric_ingress_dst_lkp` is applied, so we assume that
 283 the control plane will only call action `term_cpu_packet` when `fabric_hdr_cpu` is matched
 284 as valid. Again, our implementation prints these assumptions as warnings to the programmer,
 285 so they can confirm that the control plane will satisfy these assumptions.

286 2.2.5 Default Action Bugs

287 Finally, the *default action* errors occur when the programmer incorrectly assumes that a
 288 table performs some action when a packet misses. The NETCACHE program (described in
 289 Section 2.2.2) exhibits an example of this bug, too. The bug is shown in Figure 8, where the
 290 table `add_value_header_1` is expected to make the `nc_value_1` header valid, which is done
 291 in the `add_value_header_1_act` action. The control plane may refuse to add any rules to
 292 the table, which would cause all packets to miss, meaning that the `add_value_header_1_act`
 293 action would never be called and `nc_value_1` may not be valid. To fix this error, we simply
 294 set the default action for the table to `add_value_header_1_act`, which will force the table
 295 to remove the header no matter what the controller does.

296 2.3 A Typing Discipline to Eliminate Invalid References

297 In this paper, we propose a type system to increase the safety of P4 programs by detecting
 298 and preventing the classes of bugs defined in Section 2.2. These classes of bugs all manifest
 299 when a program attempts to access an invalid header—differentiating themselves only in
 300 their syntactic provenance. The type system that we present in the next section uses a path-
 301 sensitive analysis and occurrence typing [30], to keep track of which headers are guaranteed

<pre> /* UNSAFE */ table add_value_header_1 { actions { add_value_header_1_act; } } </pre>	<pre> /* SAFE */ table add_value_header_1 { actions { add_value_header_1_act; } default_action : add_value_header_1_act(); } </pre>
--	---

■ **Figure 8** Left: unsafe code in NETCACHE; Right: our type-safe fix.

<pre> if (ethernet.etherType == 0x0800) { apply(ipv4_table); } else if (ethernet.etherType == 0x086DD) { apply(ipv6_table); } </pre>	<pre> if (valid(ipv4)) { apply(ipv4_table); } else if (valid(ipv6)) { apply(ipv6_table); } </pre>
--	---

■ **Figure 9** Left: data-dependent header validation; Right: syntactic header validation.

to be available at any program point, and rejects programs that reference headers that might be uninitialized; thus, preventing all references to invalid headers.

Of course, in general, the problem of deciding header-validity can depend on arbitrary data, so a simple type system cannot hope to fully determine all scenarios when an instance will be valid. Indeed, programmers often use a variety of data-dependent checks to ensure safety. For instance, the control snippet shown on the left-hand side of Figure 9 will not produce undefined behavior, given a parser that chooses between parsing an `ipv4` header when `ethernet.etherType` is `0x0800`, an `ipv6` header when `ethernet.etherType` is `0x86DD`, and throws a parser error otherwise.

While this code is safe in this very specific context, it quickly becomes unsafe when ported to other contexts. For example in `switch.p4`, which performs tunneling, the egress control node copies the `inner_ethernet` header into the `ethernet`; however the `inner_ethernet` header may not be valid at the program point where the copy is performed. This behavior is left undefined [7], so targets are free to read arbitrary bits, in which case it could decide to call the `ipv4_table` despite `ipv4` being invalid!

To improve the maintainability and portability of the code, we can replace the data-dependent checks with validity checks, as illustrated by the control snippet shown on the right-hand side of Figure 9. The validity checks assert precisely the preconditions for calling each table, so that no matter what context this code snippet is called in, it is impossible for the `ipv4_table` to be called when the `ipv4` header is invalid.

In the next section, we develop a core calculus for SAFE_{P4} with a type system that eliminates references to invalid headers, encouraging programmers to replace data-dependent checks with header-validity checks.

S

3 SafeP4

This section discusses our design goals for SAFE_{P4} and the choices we made to accommodate them, and formalizes the language’s syntax, small-step semantics, and type system.

3.1 Design

Our primary design goal for SAFEP4 is to develop a core calculus that models the main features of P4₁₄ and P4₁₆, while guaranteeing that all data from packet headers is manipulated in a safe and well-defined manner. We draw inspiration from Featherweight Java [10]—i.e., we model the essential features of P4, but prune away unnecessary complexity. The result is a minimal and elegant calculus that is easy to reason about, but can still express a large number of real-world data plane programs. For instance, P4 and SAFEP4 both achieve protocol independence by allowing the programmer to specify the types of packet headers and their order in the bit stream. Similarly, SAFEP4 mimics P4’s use of tables to interface with the controller and decide, at runtime what actions to execute. Hence, we model the integral aspect of the interface between the control plane and the data plane.

So what features does SAFEP4 prune away? We omit a number of constructs that are secondary to how packets are processed—e.g., `field_list_calculations`, `parser_exceptions`, `counters`, `meters`, etc. It would be relatively straightforward to add these to the calculus—indeed, most are already handled in our prototype—at the cost of making it more complicated. We also modify or distill several aspects of P4. For instance, P4 separates the parsing phase and the control phase. Rather than unnecessarily complicating the syntax of SAFEP4, we allow the syntactic objects that represent parsers and controls to be freely mixed. We make a similar simplification in actions, informally enforcing which primitive commands can be invoked within actions (e.g., field modification, but not conditionals).

Another challenge arises in trying to model core behaviors of both P4₁₄ and P4₁₆, in that they each have different type systems and behaviors for evaluating expressions! Our calculus abstracts away expression typing and syntax variants by assuming that we are given a set of constants k that can represent values like 0 or `True`, or operators such as `&&` and `?:`. We also assume that these operators are assigned appropriate (i.e., sound) types. With these features in hand, one can instantiate our type system over arbitrary constants.

Another departure from P4 is related to `add` command, which presents a complication for our expression types. The analogous `add_header` action in P4 simply modifies the validity bit, without initializing any of the fields. This means that accessing any of the header fields before they have been manually initialized reads a non-deterministic value. Our calculus neatly sidesteps this issue by defining the semantics of the `add(h)` primitive to initialize each of the fields of h to a default value. We assume that along with our type constants there is a function `init` that accepts a header type η and produces a header instance of type η with all fields set to their default value. Note that we could have instead modified our type system to keep track of the definedness of header fields as well as their validity. However, for simplicity we choose to focus on header validity in this paper.

The portion of our type system that analyzes header validity, requires some way of keeping track of which headers are valid. Naively, we can keep track of a set of which headers are guaranteed to be valid on all program paths, and reject programs that reference headers not in this set. However, this coarse-grained approach would lead to a large number of false positives. For instance, the parser shown in Figure 1 parses an `ethernet` header and then either boots to `ingress` or parses an `ipv4` header and then either proceeds to the `ingress` or parses a `vlan` header. Hence, at the `ingress` node, the only header that is guaranteed to be valid is the `ethernet` header! However, it is certainly safe to write an `ingress` program that references the `vlan` header after checking it was valid. To reflect this in the type system we introduce a special construct called `valid(h) c1 else c2`, which executes c_1 if h is valid and c_2 otherwise. When we type check this command, following previous work on occurrence typing [30], we check c_1 with the additional fact that h is valid, and we check c_2 with the

377 additional fact that c_2 is valid.

378 Even with this enhancement, this type system would still be overly restrictive. To see
 379 why, let us augment the parser from Figure 1 with the ability to parse TCP and UDP packets:
 380 after parsing the `ipv4` header, the parser can optionally extract the `vlan`, `tcp`, or `udp` header
 381 and then boot to ingress. Now suppose that we have a table `tcp_table` that refers to both
 382 `ipv4` and `tcp` in its `reads` declaration, and that `tcp_table` is (unsafely) applied immediately
 383 in the `ingress`. Because the validity of `tcp` implies the validity of `ipv4`, it should be safe to
 384 check the validity of `tcp` and then apply `tcp_table`. However, using the representation of
 385 valid headers as a set, we would need to ascertain the validity of `ipv4` and of `tcp`.

386 To solve this problem, we enrich our type representation to keep track of dependencies
 387 between headers. More specifically, rather than representing all headers guaranteed to be
 388 valid in a set, we use a finer-grained representation—a set of sets of headers that might be
 389 valid at the current program point. For a given header reference to be safe, it must be
 390 a member of all possible sets of headers—i.e., it must be valid on all paths through the
 391 program that reach the reference.

392 Overall, the combination of an expressive language of types and a simple version of
 393 occurrence typing allows us to capture dependencies between headers and perform useful
 394 static analysis of the dynamic property of header validity.

395 The final challenge with formally modelling P4 is its interface with the control-plane,
 396 which populates the tables and provides arguments to the actions. While the control-plane
 397 only methodology for managing switch behavior is to populate the match-action tables with
 398 forwarding entries, it is perfectly capable of producing undefined behavior. However, if we
 399 assume that the controller is well-intentioned, we can prove the safety of more programs.

400 In our formalization, to streamline the presentation, we model the control plane as a
 401 function $\mathcal{CA}(t, H) = (a_i, \bar{v})$ that takes in a table t and the current headers H and produces
 402 the action to call a_i and the (possibly empty) action data arguments \bar{v} . We also use a
 403 function $\mathcal{CV}(t) = (\bar{S}, \bar{e})$ that analyzes a table t and produces a list of match key expressions
 404 \bar{e} that must be evaluated when the table is invoked, and a list of sets of valid headers \bar{S} , one
 405 set for each action, that can be safely assumed valid when the entries are populated by the
 406 control plane. Together, these functions model the runtime interface between the switch
 407 and the controller. In order to prove progress and preservation, we assume that \mathcal{CV} and \mathcal{CA}
 408 satisfy three simple correctness properties—see Appendix C.1 for details.

409 3.2 Syntax

410 The syntax of SAFE P4 is shown in Figure 10. To lighten the notation, we write \bar{x} as
 411 shorthand for a (possibly empty) sequence x_1, \dots, x_n .

412 A SAFE P4 program consists of a sequence of declarations \bar{d} and a command c . The set
 413 of declarations includes header types, header instances, and tables. Header type declarations
 414 describe the format of individual headers and are defined in terms of a name and a sequence
 415 of field declarations. The notation “ $f : \tau$ ” indicates that field f has type τ . We let η range
 416 over header types. A header instance declaration assigns a name h to a header type η . The
 417 map \mathcal{HT} encodes the (global) mapping between header instances and header types. Table
 418 declarations $t(\bar{e}, \bar{a})$, are defined in terms of a sequence of match-key expressions \bar{e} read in
 419 the table, and a sequence of actions \bar{a} . The notation $t.reads$ denotes the expressions and
 420 $t.actions$ denotes the actions.

421 Actions are written as (uncurried) λ -abstractions. An action $\lambda \bar{x}. c$ declares a (possibly
 422 empty) sequence of parameters, drawn from a fresh set of names, which are in scope for the
 423 command c . The run-time arguments for actions are chosen by the control plane. Note that

<div>Commands</div> <div><div><div><div><div>$c ::=$</div></div><div><div><div><div>$$</div><div>$extract(h)$</div></div><div><div><div>EXTRACTION</div></div></div></div><div><div><div>$$</div><div>$emit(h)$</div></div><div><div><div>DEPARSING</div></div></div></div><div><div><div>$$</div><div>$c_1; c_2$</div></div><div><div><div>SEQUENCE*</div></div></div></div><div><div><div>$$</div><div>$if(e) \ c_1 \ else \ c_2$</div></div><div><div><div>CONDITIONAL</div></div></div></div><div><div><div>$$</div><div>$valid(h) \ c_1 \ else \ c_2$</div></div><div><div><div>VALIDITY</div></div></div></div><div><div><div>$$</div><div>$t.apply()$</div></div><div><div><div>APPLICATION</div></div></div></div><div><div><div>$$</div><div>$skip$</div></div><div><div><div>SKIP</div></div></div></div><div><div><div>$$</div><div>$add(h)$</div></div><div><div><div>ADDITION*</div></div></div></div><div><div><div>$$</div><div>$remove(h)$</div></div><div><div><div>REMOVAL*</div></div></div></div><div><div><div>$$</div><div>$h.f = e$</div></div><div><div><div>MODIFICATION*</div></div></div></div></div></div></div></div>	<div><div>Declarations</div><div><div><div><div><div>$d ::=$</div></div><div><div><div><div>$$</div><div>$t(\bar{e}, \bar{a})$</div></div><div><div><div>TABLE</div></div></div></div><div><div><div>$$</div><div>$\eta \ \{f : \tau\}$</div></div><div><div><div>HEADER TYPE</div></div></div></div><div><div><div>$$</div><div>$h \mapsto \eta$</div></div><div><div><div>INSTANTIATION</div></div></div></div></div></div></div></div><div><div>Program</div><div><div><div><div>$\mathcal{P} ::=$</div></div><div><div><div>(\bar{d}, c)</div></div><div><div><div>PROGRAM</div></div></div></div></div></div></div></div>
<div><div>Actions</div><div><div><div><div>$a ::=$</div></div><div><div><div>$\lambda \bar{x}. c$</div></div><div><div><div>ACTION</div></div></div></div></div></div></div>	
<div><div>Expressions</div><div><div><div><div><div>$e ::=$</div></div><div><div><div><div>$$</div><div>v</div></div><div><div><div>VALUES</div></div></div></div><div><div><div>$$</div><div>$h.f$</div></div><div><div><div>HEADER FIELD</div></div></div></div><div><div><div>$$</div><div>x</div></div><div><div><div>VARIABLE</div></div></div></div><div><div><div>$$</div><div>k^n</div></div><div><div><div>CONSTANT</div></div></div></div></div></div></div></div></div>	<div><div><div>Action Types</div><div>Expression Types</div></div><div><div><div><div>$\alpha ::=$</div></div><div><div><div>$\bar{\tau} \rightarrow \Theta$</div></div><div><div><div>$\tau ::=$</div></div><div><div><div>Bool</div></div></div></div><div><div><div>$$</div><div>$\bar{\tau} \rightarrow \tau$</div></div><div><div><div>$$</div><div>\dots</div></div></div></div></div></div></div></div>

■ **Figure 10** Syntax of SAFE4

we artificially restrict the commands that can be called in the body of the action to addition, removal, modification and sequence; these actions are identified with an asterisk in Figure 10.

The calculus provides commands for extracting (*extract*), creating (*add*), removing (*remove*), and modifying ($h.f = e$) header instances. The *emit* command serializes a packet header back into a bit sequence (*emit*). The *if*-statement conditionally executes one of two commands based on the value of a boolean condition. Similarly, the *valid*-statement branches on the validity of h . Table application commands ($t.apply()$) are used to invoke a table t in the current state. The *skip* command is a no-op.

The only built-in expressions in SAFE4 are variables x and header fields, written $h.f$. We let v range over values and assume a collection of n -ary constant operators $k^n \in K$.

For simplicity, we assume that every header referenced in an expression has a corresponding instance declaration. We also assume that header instance names h , header type names η , variable names x , and table names t are drawn from disjoint sets of names H, E, V , and T respectively and that each name is declared only once.

3.3 Type System

SAFE4 provides two main kinds of types, basic types τ and header types Θ as shown in Figure 10. We assume that the set of basic types includes booleans (for conditionals) as well as tuples and function types (for actions).

A header type Θ represents a set of possible co-valid header instances. The type 0 denotes

$\begin{aligned} \llbracket 0 \rrbracket &= \{\} \\ \llbracket 1 \rrbracket &= \{\{\}\} \\ \llbracket h \rrbracket &= \{\{h\}\} \\ \llbracket \Theta_1 \cdot \Theta_2 \rrbracket &= \llbracket \Theta_1 \rrbracket \bullet \llbracket \Theta_2 \rrbracket \\ \llbracket \Theta_1 + \Theta_2 \rrbracket &= \llbracket \Theta_1 \rrbracket \cup \llbracket \Theta_2 \rrbracket \end{aligned}$	$\begin{aligned} \mathcal{F}(h, f_i) &= \tau_i & \text{Field lookup} \\ \mathcal{A}(a) &= \lambda \bar{x} : \bar{\tau}. c & \text{Action lookup} \\ \mathcal{CA}(H, t) &= (a_i, \bar{v}) & \text{Control-plane actions} \\ \mathcal{CV}(t) &= (\bar{S}, \bar{e}) & \text{Control-plane validity} \end{aligned}$
---	---

■ **Figure 11** Semantics of header types (left) and auxiliary functions (right).

the empty set. This type arises when there are unsatisfiable assumptions about which headers are valid. The type 1 denotes the singleton denoting the empty set of headers. It describes the type of the initial state of the program. The type h denotes a singleton set, $\{\{h\}\}$ —i.e., states where only h is valid. The type $\Theta_1 \cdot \Theta_2$ denotes the set obtained by combining headers from Θ_1 and Θ_2 —i.e., a product or concatenation. Finally, the type $\Theta_1 + \Theta_2$ denotes the union of Θ_1 or Θ_2 , which intuitively, represents an alternative.

The semantics of header types, $\llbracket \Theta \rrbracket$, is defined by the equations in Figure 11. Intuitively, each subset represents one alternative set of headers that may be valid. For example, the header type `eth · (ipv4 + 1)` denotes the set $\{\{\text{eth}, \text{ipv4}\}, \{\text{eth}\}\}$.

To formulate the typing rules for SAFE_{P4}, we also define a set of operations on header types: **Restrict**, **NegRestrict**, **Includes**, **Remove**, and **Empty**. The restrict operator **Restrict** Θh recursively traverses Θ and keeps only those choices in which h is contained, mapping all others to 0. Semantically this has the effect of throwing out the subsets of $\llbracket \Theta \rrbracket$ that do not contain h . Dually **NegRestrict** Θh produces only those choices/subsets where h is invalid. **Includes** Θh traverses Θ and checks that h is always valid. Semantically this says that h is a member of every element of $\llbracket \Theta \rrbracket$. **Remove** Θh removes h from every path, which means, semantically that it removes h from every element of $\llbracket \Theta \rrbracket$. Finally, **Empty** Θ checks whether Θ denotes the empty set. An in-depth treatment of these operators, with proofs of all of these claims can be found in Appendix B.

3.3.1 Typing Judgement

The main typing judgement has the form $\Gamma \vdash c : \Theta \Rightarrow \Theta'$, which means that in variable context Γ , if c is executed in the header context Θ , then a header instance type Θ' is assigned. Intuitively, Θ encodes the sets of headers that may be valid when type checking a command. Γ is a standard type environment which maps variables x to type τ . If there exists Θ' such that $\Gamma \vdash c : \Theta \Rightarrow \Theta'$, we say that c is well-typed in Θ .

The typing rules rely on several auxiliary definitions shown in Figure 11. The field type lookup function $\mathcal{F}(h, f_i)$ returns the type assigned to a field f_i in header h . The action lookup function $\mathcal{A}(a)$ returns the action definition $\lambda \bar{x} : \bar{\tau}. c$ for action a . Finally, the function $\mathcal{CA}(t, H)$ computes the run-time actions for table t , while $\mathcal{CV}(t)$ computes t 's assumptions about validity. Both of these are assumed to be instantiated by the control plane in a way that satisfies basic correctness properties—see Appendix C.1.

The typing rules for commands are presented in Figure 12. The rule T-ZERO gives a command an arbitrary output type if the input type is empty. The rules T-SKIP are T-SEQ are standard. The rule T-IF a path-sensitive union type between the type computed for each branch. The rule T-IFVALID is similar, but leverages knowledge about the validity of

$\frac{\text{T-ZERO} \quad \text{Empty } \Theta_1}{\Gamma \vdash c : \Theta_1 \Rightarrow \Theta_2}$ $\frac{\text{T-SEQ} \quad \Gamma \vdash c_1 : \Theta \Rightarrow \Theta_1 \quad \Gamma \vdash c_2 : \Theta_1 \Rightarrow \Theta_2}{\Gamma \vdash c_1; c_2 : \Theta \Rightarrow \Theta_2}$ $\frac{\text{T-IF} \quad \Gamma; \Theta \vdash e : \text{Bool} \quad \Gamma \vdash c_1 : \Theta \Rightarrow \Theta_1 \quad \Gamma \vdash c_2 : \Theta \Rightarrow \Theta_2}{\Gamma \vdash \text{if } (e) \ c_1 \ \text{else } c_2 : \Theta \Rightarrow \Theta_1 + \Theta_2}$ $\frac{\text{T-IFVALID} \quad \Gamma \vdash c_1 : \text{Restrict } \Theta \ h \Rightarrow \Theta_1 \quad \Gamma \vdash c_2 : \text{NegRestrict } \Theta \ h \Rightarrow \Theta_2}{\Gamma \vdash \text{valid}(h) \ c_1 \ \text{else } c_2 : \Theta \Rightarrow \Theta_1 + \Theta_2}$ $\frac{\text{T-MOD} \quad \text{Includes } \Theta \ h \quad \mathcal{F}(h, f) = \tau_i \quad \Gamma; \Theta \vdash e : \tau_i}{\Gamma \vdash h.f = e : \Theta \Rightarrow \Theta}$	$\frac{\text{T-EXTR}}{\Gamma \vdash \text{extract}(h) : \Theta \Rightarrow \Theta \cdot h}$ $\frac{\text{T-EMIT}}{\Gamma \vdash \text{emit}(h) : \Theta \Rightarrow \Theta}$ $\frac{\text{T-ADD}}{\Gamma \vdash \text{add}(h) : \Theta \Rightarrow \Theta \cdot h}$ $\frac{\text{T-REM}}{\Gamma \vdash \text{remove}(h) : \Theta \Rightarrow \text{Remove } \Theta \ h}$ $\frac{\text{T-APPLY} \quad \mathcal{CV}(t) = (\bar{S}, \bar{e}) \quad t.\text{actions} = \bar{a} \quad \cdot; \Theta \vdash e_i : \tau_i \text{ for } e_i \in \bar{e} \quad \text{Restrict } \Theta \ S_i \vdash a_i : \bar{\tau}_i \rightarrow \Theta'_i \text{ for } a_i \in \bar{a}}{\Gamma \vdash t.\text{apply}() : \Theta \Rightarrow \left(\sum \Theta'_i \right)}$
--	---

■ **Figure 12** Command typing Rules for SAFEP4

$\frac{\Gamma, \bar{x} : \bar{\tau} \vdash c : \Theta \Rightarrow \Theta'}{\Gamma; \Theta \vdash \lambda \bar{x} : \bar{\tau}. c : \bar{\tau} \rightarrow \Theta'} \quad (\text{T-ACTION})$

■ **Figure 13** Action typing rule for SAFEP4

h . So the true branch c_1 is checked in the context **Restrict** $\Theta \ h$, and the false branch c_2 is checked in the context **NegRestrict** $\Theta \ h$. The top-level output type is the union of the resulting output types for c_1 and c_2 . The rule T-MOD checks that h is guaranteed to be valid using the **Includes** operator, and uses the auxiliary function \mathcal{F} to obtain the type assigned to $h.f$. Note that the set of valid headers does not change when evaluating an assignment, so the output and input types are identical. The rules T-EXTR and T-ADD assign header extractions and header additions the type $\Theta \cdot h$, reflecting the fact that h is valid after the command executes. Emitting packet headers does not change the set of valid headers, which is captured by rule T-EMIT. The typing rule T-REM uses the **Remove** operator to remove h from the input type Θ . Finally, the rule T-APPLY checks table applications. To understand how it works, let us first consider a simpler but less precise typing rule:

$$\frac{t.\text{reads} = \bar{e} \quad \cdot; \Theta \vdash e_i : \tau_i \text{ for } e_i \in \bar{e} \quad t.\text{actions} = \bar{a} \quad \cdot; \Theta \vdash a_i : \bar{\tau}_i \rightarrow \Theta'_i \text{ for } a_i \in \bar{a}}{\cdot \vdash t.\text{apply}() : \Theta \Rightarrow \left(\sum \Theta'_i \right)}$$

474 Intuitively, this rule says that to type check an table application, we check each expression it
 475 reads and each of its actions. The final header type is the union of the types computed for

$\frac{\text{T-CONST} \quad \text{typeof}(k) = \bar{\tau} \rightarrow \tau' \quad \Gamma; \Theta \vdash e_i : \tau_i}{\Gamma; \Theta \vdash k(\bar{e}) : \tau'}$	$\frac{\text{T-VAR} \quad x : \tau \in \Gamma}{\Gamma; \Theta \vdash x : \tau}$	$\frac{\text{T-FIELD} \quad \text{Includes } \Theta \ h \quad \mathcal{F}(h, f) = \tau}{\Gamma; \Theta \vdash h.f : \tau} \quad (\text{T-FIELD})$
--	---	---

■ **Figure 14** Expression typing rules for SAFEP4

the actions. To put it another way, it models table application as a non-deterministic choice between its actions. However, while this rule is sound, it is overly conservative. In particular, it does not model the fact that the control plane often uses header validity bits to control which actions are executed. Hence, The actual typing rule, T-APPLY, is parameterized on a function $\mathcal{CV}(t)$ that models the choices made by the control plane, returning for each action a_i , a set of headers S_i that can be assumed valid when type checking a_i , as well as a subset of the expressions read by the table — e.g., excluding expressions that can be wildcarded when certain validity bits are false.

The typing judgement for actions (Figure 13) is of the form $\Gamma; \Theta \vdash a : \bar{\tau} \rightarrow \Theta$, meaning that a has type $\bar{\tau} \rightarrow \Theta$ in variable context Γ and header context Θ . Given a variable context Γ and header type Θ , an action $\lambda \bar{x}.c$ encodes a function of type $\bar{\tau} \rightarrow \Theta'$, so long as the body c is well-typed in the context where Γ is extended with $x_i : \tau_i$ for every i .

The typing rules for expressions are shown in Figure 14. Constants are typechecked according to rule T-CONSTANT, as long as each expression that is passed as an argument to the constant k has the type required by the `typeof` function. The rule T-VAR is standard.

3.4 Operational Semantics

We now present the small-step operational semantics of SAFEP4. We define the operational semantics for commands in terms of four-tuples $\langle I, O, H, c \rangle$, where I is the input bit stream (which is assumed to be infinite for simplicity), O is the output bit stream, H is a map that associates each valid header instance with a records containing the values of each field, and c is the command to be evaluated. The reduction rules are presented in Figure 15.

The command $\text{extract}(h)$ evaluates via the rule E-EXTR, which looks up the header type in \mathcal{HT} and then invokes corresponding deserialization function. The deserialized header value v is added to the map of valid header instances, H . For example, assuming the header type $\eta = \{f : \text{bit}\langle 3 \rangle; g : \text{bit}\langle 2 \rangle\}$ has two fields f and g and $I = 11000B$ where B is the rest of the bit stream following, then $\text{deserialize}_\eta(I) = (\{f = 110; g = 00\}, B)$.

The rule E-EMIT serializes a header instance h back into a bit stream. It first looks up the corresponding header type and header value in the header table \mathcal{HT} and the map of valid headers respectively. The header value is then passed to the serialization function for the header type to produce a bit sequence that is appended to the output bit stream. Similarly, we assume that a serialization function is defined for every header type, which takes the bit values of the fields of a header value and concatenates them to produce a single bit sequence. We adopt the semantics of P4 with respect to emitting invalid headers. Emitting an invalid header instance, i.e., a header instance which has not been added or extracted, has just a no-op without any effect on the output bit stream (rule E-EMITINVALID). Notice also that the header remains unchanged in H .

Sequential composition reduces left to right, i.e., the left command needs to be reduced to *skip* before the right command can be reduced (rule E-SEQ). The evaluation of conditionals (rules E-IF, E-IFTRUE, E-IFFALSE) is standard. Both E-SEQ and E-IF are relegated to the

$\frac{\text{E-EXTR} \quad \mathcal{HT}(h) = \eta \quad \text{deserialize}_\eta(I) = (v, I')}{\langle I, O, H, \text{extract}(h) \rangle \rightarrow \langle I', O, H[h \mapsto v], \text{skip} \rangle}$	$\frac{\text{E-MOD} \quad H(h) = r \quad r' = \{r \text{ with } f = v\}}{\langle I, O, H, h.f = v \rangle \rightarrow \langle I, O, H[h \mapsto r'], \text{skip} \rangle}$
$\frac{\text{E-EMIT} \quad \mathcal{HT}(h) = \eta \quad \text{serialize}_\eta(H(h)) = \bar{B}}{\langle I, O, H, \text{emit}(h) \rangle \rightarrow \langle I, O, \bar{B}, H, \text{skip} \rangle}$	$\frac{\text{E-APPLY} \quad \mathcal{CA}(H, t) = (a_i, \bar{v}) \quad \mathcal{A}(a_i) = \lambda \bar{x}.c_i}{\langle I, O, H, t.\text{apply}() \rangle \rightarrow \langle I, O, H, c_i[\bar{v}/\bar{x}] \rangle}$
$\frac{\text{E-EMITINVALID} \quad h \notin \text{dom}(H)}{\langle I, O, H, \text{emit}(h) \rangle \rightarrow \langle I, O, H, \text{skip} \rangle}$	$\frac{\text{E-ADD} \quad \mathcal{HT}(h) = \eta \quad \text{init}_\eta = v}{\langle I, O, H, \text{add}(h) \rangle \rightarrow \langle I, O, H[h \mapsto v], \text{skip} \rangle}$
$\frac{\text{E-IFVALIDTRUE} \quad h \in \text{dom}(H)}{\langle I, O, H, \text{valid}(h) \ c_1 \ \text{else} \ c_2 \rangle \rightarrow \langle I, O, H, c_1 \rangle}$	$\frac{\text{E-ADDVALID} \quad h \in \text{dom}(H)}{\langle I, O, H, \text{add}(h) \rangle \rightarrow \langle I, O, H, \text{skip} \rangle}$
$\frac{\text{E-IFVALIDFALSE} \quad h \notin \text{dom}(H)}{\langle I, O, H, \text{valid}(h) \ c_1 \ \text{else} \ c_2 \rangle \rightarrow \langle I, O, H, c_2 \rangle}$	$\frac{\text{E-REM}}{\langle I, O, H, \text{remove}(h) \rangle \rightarrow \langle I, O, H \setminus h, \text{skip} \rangle}$

■ **Figure 15** Selected rules of the operational semantics of SAFEP4; the elided rules are standard and can be found in Appendix A.

$\frac{\text{E-CONST} \quad \llbracket k \rrbracket(v_1, \dots, v_n) = v}{\langle H, k(v_1, \dots, v_n) \rangle \rightarrow v}$	$\frac{\text{E-FIELD} \quad H(h) = \{f_1 : n_1, \dots, f_k : n_k\}}{\langle H, h.f_i \rangle \rightarrow n_i}$
---	--

■ **Figure 16** Selected rules of the operational semantics for expressions.

515 appendix for brevity. The rules for validity checks (E-IFVALIDTRUE, E-IFVALIDFALSE) step
 516 to the true branch if $h \in \text{dom}(H)$ and to the false branch otherwise.

517 Table application commands are evaluated according to rule E-TAPPLY. We first invoke
 518 the control plane function $\mathcal{CA}(H, t)$, which determines an action a_i and action data v . Then
 519 we use \mathcal{A} to lookup the definition of a_i , yielding $\lambda \bar{x} : \bar{\tau}. c_i$ and step to $c_i[\bar{v}/\bar{x}]$. Note that for
 520 simplicity, we model the evaluation of expressions read by the table in the the control-plane
 521 function \mathcal{CA} rather than in the calculus.

522 The rule E-ADD evaluates addition commands $\text{add}(h)$. Similar to header extraction, the
 523 $\text{init}_\eta()$ function produces a header instance v of type η with all fields set to a default value
 524 and extends the map H with $h \mapsto v$. Note that according to E-ADD-EXIST, if the header
 525 instance is already valid, $\text{add}(h)$ does nothing. Finally, the rule E-REM removes the header
 526 from the map H . Again, if a header h is already invalid, removing it has no effect.

527 The semantics for expressions is given in Figure 16. We assume that there is an evaluation
 528 function for constants $\llbracket k \rrbracket(\bar{v}) = v$ that is well-behaved—i.e., if $\text{typeof}(k) = \bar{\tau} \rightarrow \tau'$ and $\bar{v} : \bar{\tau}$,
 529 then $.; \vdash \llbracket k \rrbracket(\bar{v}) : \tau'$. We use these facts to prove progress and preservation.

530 We define the semantics of expressions using tuples $\langle H, e \rangle$, where H is the same map
 531 used in the semantics of commands and e is the expression to evaluate. The rule E-CONST
 532 evaluates constants (we omit the obvious congruence rule) and rule E-FIELD reduces header
 533 field expressions to the value stored in H for the respective field.

ENT-SEQ				
ENT-EMPTY	ENT-INST	$H_1 \models \Theta_1$	ENT-CHOICE L	ENT-CHOICE R
$\frac{}{\cdot \models 1}$	$\frac{\text{dom}(H) = \{h\}}{H \models h}$	$\frac{H_2 \models \Theta_2}{H_1 \cup H_2 \models \Theta_1 \cdot \Theta_2}$	$\frac{H \models \Theta_1}{H \models \Theta_1 + \Theta_2}$	$\frac{H \models \Theta_2}{H \models \Theta_1 + \Theta_2}$

■ **Figure 17** The *Entailment* relation between header instances and header instance types

3.5 Safety of SafeP4

We prove safety in terms of progress and preservation. Both theorems make use of the relation $H \models \Theta$ which intuitively holds if H is described by Θ . The formal definition, as given in Figure 17, satisfies $H \models \Theta$ if and only if $\text{dom}(H) \in \llbracket \Theta \rrbracket$.

We prove type safety via progress and preservation theorems. The respective proofs are mostly straightforward for our system—we highlight the unusual and nontrivial cases below and relegate the full proofs to the appendix.

► **Theorem 1 (Progress).** *If $\Gamma \vdash c : \Theta \Rightarrow \Theta'$ and $H \models \Theta$, then either,*

- $c = \text{skip}$, or
- $\exists \langle I', O', H', c' \rangle. \langle I, O, H, c \rangle \rightarrow \langle I', O', H', c' \rangle$.

Intuitively, progress says that a well-typed command is fully reduced or can take a step.

► **Theorem 2 (Preservation).** *If $\Gamma \vdash c : \Theta_1 \Rightarrow \Theta_2$ and $\langle I, O, H, c \rangle \rightarrow \langle I', O', H', c' \rangle$, where $H \models \Theta_1$, then $\exists \Theta'_1, \Theta'_2. \Gamma \vdash c : \Theta'_1 \Rightarrow \Theta'_2$ where $H' \models \Theta'_1$ and $\Theta'_2 < \Theta_2$.*

More interestingly, preservation says that if a command c is well-typed with input type Θ_1 and output type Θ_2 , and c evaluates to c' in a single step, then there exists an input type Θ'_1 and an output type Θ'_2 that make c' well-typed. To make the inductive proof go through, we also need to prove that Θ'_1 describes the same maps of header instance H as Θ_1 , and Θ'_2 is semantically contained in Θ_2 . (These conditions are somewhat reminiscent of conditions found in languages with subtyping.)

Proof. By induction on a derivation of $\Gamma \vdash c : \Theta_1 \Rightarrow \Theta_2$, with a case analysis on the last rule used. We focus on two of the most interesting cases. See Appendix C for the full proof.

Case T-IFVALID: $c = \text{valid}(h) \ c_1 \ \text{else} \ c_2$ and $\Gamma \vdash c_1 : \text{Restrict } \Theta_1 \ h \Rightarrow \Theta_{12}$ and $\Gamma \vdash c_2 : \text{NegRestrict } \Theta_1 \ h \Rightarrow \Theta_{22}$ and $\Theta_2 = \Theta_{12} + \Theta_{22}$.

There are two evaluation rules that apply to c , E-IFVALIDTRUE and E-IFVALIDFALSE

Subcase E-IFVALIDTRUE: $c' = c_1$ and $h \in \text{dom}(H)$ and $H' = H$.


Let $\Theta'_1 = \text{Restrict } \Theta_1 \ h$ and $\Theta'_2 = \Theta_{12}$. We have $\Gamma \vdash c' : \Theta'_1 \Rightarrow \Theta'_2$ by assumption, we have $H \models \Theta'_1$ by Lemma 17, and we have $\Theta'_2 < \Theta_2$ by the definition of $<$ and the semantics of union.

Subcase E-IFVALIDFALSE: $c' = c_2$ and $h \notin \text{dom}(H)$ and $H' = H$.


Symmetric to the previous case.

Case T-APPLY: $c = t.\text{apply}()$ and $\mathcal{CV}(t) = (\bar{S}, \bar{e})$ and $t.\text{actions} = \bar{a}$ and $\vdash \Theta \vdash e_i : \tau_i$ for $e_i \in \bar{e}$ and $\text{Restrict } \Theta_1 \ S_i \vdash a_i : \bar{\tau}_i \rightarrow \Theta'_i$ for $a_i \in \bar{a}$ and $\Theta_2 = \sum (\Theta'_i)$

Only one evaluation rule that applies to c , E-APPLY. It follows that $\mathcal{CA}(H, t) = (a_i, \bar{v})$, and $c' = c_i[\bar{v}/\bar{x}]$ where $\mathcal{A}(a_i) = \lambda \bar{x}. c_i$. By inverting T-Action, we have $\Gamma, \bar{x} : \bar{\tau}_i \vdash c_i :$



■ **Figure 18** Proportional frequencies of each bug type per-program. The raw number of bugs for each program and category is reported at the top of each stacked bar.



■ **Figure 19** Frequency of each bug across all programs. The raw number of bugs in each category is reported to the right of the bar

568 **Restrict** $\Theta S_i \Rightarrow \Theta'_i$. By Proposition 14, we have $\cdot; \cdot \vdash \bar{v} : \bar{\tau}_i$. By the substitution lemma,
 569 we have $\Gamma \vdash c_i[\bar{w}/\bar{x}] : \text{Restrict } \Theta S_i \Rightarrow \Theta'_i$. Let $\Theta'_1 = \text{Restrict } \Theta S_i$ and $\Theta'_2 = \Theta'_i$. We
 570 have shown that $\Gamma \vdash c' : \Theta'_1 \Rightarrow \Theta'_2$, we have that $H' \models \Theta'_1$ by Proposition 15, and we have
 571 $\Theta'_2 < \Theta_2$ by the definition of $<$ and the semantics of union types. ◀

572 4 Experience (Evaluation)

573 We implemented our type system in a tool called P4CHECK that automatically checks P4
 574 programs and reports violations of the type system presented in Figure 12. P4CHECK uses
 575 the front-end of p4v [18] and handles the full P4₁₄ language. Our key findings, which are
 576 reported in detail below, show (i) that our type system finds bugs “in the wild” and (ii) that
 577 the programmer effort needed to repair programs to pass our type checker is modest.

578 4.1 Overview of Bugs in the Wild

579 We ran P4CHECK on 15 open source P4₁₄ programs³ designed for research and industrial
 580 use. The subject programs are of varying sizes and complexities—ranging from 143 to 9060
 581 lines of code. Our criteria for selecting programs was: (1) each program had to be open

³ We chose to check P4₁₄ instead of P4₁₆, since there were more realistic open-source programs available on GitHub in P4₁₄. In fact, we could find no P4₁₆ program that comes close to the size and complexity of `switch.p4`, which is written in P4₁₄.

source, (2) available on GitHub, and (3) compile without errors, (4) and be written either by industrial teams developing production code or by researchers implementing standard or novel network functionality in P4 (i.e., we excluded programs primarily used for teaching). Out of the 15 subject programs only 4 passed our type checker, all of which were simple implementations of routers or DDoS mitigation that accepted only a small number of packet types and were relatively small (188 - 635 lines of code). For the remaining 11 programs (industrial and research) our checker found 418 type checking violations overall.

Frequently, multiple violations produced by P4CHECK have the same bug as their root cause. For example, if a single action `rewrite_ipv4` that rewrites fields `srcAddr` and `dstAddr` for an `ipv4` header is called in a context that cannot prove that `ipv4` is valid, then both references to `ipv4.srcAddr` and `ipv4.dstAddr` will be reported as violations, even though they are due to the same *control* bug (Section 2.2.2)—namely that `rewrite_ipv4` was not called in a context that could prove the validity of `ipv4`. To address this issue, we applied another metric to quantify the number of bugs (inspired by the method proposed by others [15]): we equate the number of bugs in each program with the number of bug *fixes* required to make the program in question pass our type checker. Using this metric, we counted 58 bugs.

We classified the bugs according to the classes described in Section 2.2. Figure 18 depicts the per-program breakdown of the frequency of each bug class, and Figure 19 depicts the overall frequency of each bug. Notice that even though table action bugs were the most frequent bug (with 22 occurrences), they were only found in a single program (`switch.p4`). These bugs are especially prevalent in this program because of its heavy reliance on correct control-plane configuration. Conversely, there were 9 occurrences across 5 programs for both parser bugs and table reads bugs.

Readers familiar with previous work on `p4v` [18], a recent P4 verification tool may notice that we detected no default action bugs for the `switch.p4` program, while `p4v` reported many! The reasons for this are twofold. First, `p4v` allows programmers to verify arbitrarily complex propositions in Dijkstra’s guarded command language, which means that it can express fine-grained conditions on tables and relationships between them. In contrast, we make heuristic assumptions about P4 programs that automatically eliminate many bugs, including some default action bugs. Second, our repairs are often coarse-grained and may enforce a stronger guarantee on the program than may be necessary; using first-order logic annotations, `p4v` programmers manually specify the most liberal (and hence complex) assumptions.

We make no claims about the completeness of our taxonomy. For example, we found one instance, in the HAPPYFLOWFRIENDS program, where the programmer had mistakenly instantiated metadata *m* as a header, and consequently did not parse *m* (since metadata is always valid) causing *m* to (ironically) always be invalid.

4.2 P4Check in Action

We reprise the canonical examples of each class of bugs from Section 2.2, describing how P4CHECK detects them and discussing ways to fix them.

4.2.1 Parser Bugfixes

Recall Figure 4, which exhibits the parser bug. The bug occurs because the parser, which extracts IPv4-TCP packets, boots unexpected packets (such as IPv6 or UDP packets) directly to `ingress`, which then assumes that both the `ipv4` and `tcp` headers are valid, even though the parser does not guarantee this fact.

```

./h.p4, line 350, cols 12-21: error tcp not guaranteed to be valid
./h.p4, line 118, cols 8-16: error ipv4 not guaranteed to be valid
./h.p4, line 101, cols 42-50: error ipv4 not guaranteed to be valid
./h.p4, line 320, cols 8-15: error tcp not guaranteed to be valid
./h.p4, line 362, cols 12-19: error tcp not guaranteed to be valid
./h.p4, line 362, cols 29-36: error tcp not guaranteed to be valid
./h.p4, line 295, cols 60-69: error tcp not guaranteed to be valid
./h.p4, line 107, cols 8-16: error ipv4 not guaranteed to be valid
./h.p4, line 101, cols 42-50: error ipv4 not guaranteed to be valid
./h.p4, line 163, cols 8-16: error ipv4 not guaranteed to be valid
./h.p4, line 101, cols 42-50: error ipv4 not guaranteed to be valid

./h.p4, line 350, cols 12-21: error tcp not guaranteed to be valid
./h.p4, line 320, cols 8-15: error tcp not guaranteed to be valid
./h.p4, line 362, cols 12-19: error tcp not guaranteed to be valid
./h.p4, line 362, cols 29-36: error tcp not guaranteed to be valid
./h.p4, line 295, cols 60-69: error tcp not guaranteed to be valid

```

■ **Figure 20** Curated output from P4CHECK for the parser bug in NETHCF before (above) and after (below) modifying `parse_ethernet`

In terms of our type system, the parser produces packets of type `ethernet · (1 + ipv4 · (1 + tcp))`; however the control only handles packets of type `ethernet · ipv4 · tcp`. Hence, when typecheck this example, P4CHECK reports every reference to `tcp` and `ipv4` in the whole program as a violation of the type system. As shown in the top half Figure 20, we get an error message at every reference to `ipv4` or `tcp`. The ubiquity of the reports intimates a mismatch between the parsing and the control types, which gives the programmer a hint as how to fix the problem.

When we modify the `default` clause in `parse_ethernet`, as in Figure 4, and run our tool again, all of the `ipv4` violations are removed, as shown in the bottom half of Figure 20. Then fixing the `parse_ipv4` parser, as in Figure 4, causes our tool to output no violations; we count each of these fixes to be a separate bug. Now, the type on entering the `ingress` control function is `ethernet · ipv4 · tcp`, so all subsequent references to `ipv4` and `tcp` will be safe.

4.2.2 Control Bugfixes

The control bug occurs when the incoming type presents a choice between two headers that is not handled by subsequent code. The exemplar presented in Figure 5 presents a parser that produces the following type Θ :

$$\Theta = \text{ethernet} \cdot (1 + \text{ipv4} \cdot (1 + \text{udp} \cdot (1 + \text{nc_hdr} \cdot \tau) + \text{tcp})),$$

where τ is a widely branching type representing caching operations. Notice `Includes Θ nc_hdr` is false. However, the control nodes `process_cache` and `process_value` only type check in contexts where `Includes Θ nc_hdr` is true. P4CHECK reports type violations at every reference to `nc_hdr`. Fixing this error is simply a matter of wrapping the `process_cache()` call in a validity check as demonstrated in Figure 5. Since NETCACHE handles TCP and UDP packets as well as its special-purpose packets, we simply continue on and apply the IPv4 routing table if the validity check for `nc_hdr` fails.⁴

⁴ Astute readers may detect a parser bug in this example. Hint, the `ipv4_route` table requires `Includes Θ ipv4` where Θ is type where it is applied.

```

port.p4, line 248, cols 8-24: warning: assuming either vlan_tag_[0]
    matched as valid or vlan_tag_[0].vid wildcarded

port.p4, line 250, cols 8-24: warning: assuming either vlan_tag_[1]
    matched as valid or vlan_tag_[1].vid wildcarded

-----

fabric.p4 line 42, cols 41-67: warning: assuming fabric_header_cpu
    matched as valid for rules with action terminate_cpu_packet

fabric.p4, line 57, cols 17-54: warning: assuming fabric_header_unicast
    matched as valid for rules with action
    terminate_fabric_unicast_packet

fabric.p4, line 81, cols 17-56: warning: assuming
    fabric_header_multicast matched as valid for rules with action
    terminate_fabric_multicast_packet

```

■ **Figure 21** Warnings printed after fixing `switch.p4`'s reads bug (top), and its actions bug (bottom)

652 4.2.3 Table Reads Bugfixes

653 Table reads errors, like the one in Figure 6, occur when a header h is included in the `reads`
654 declaration of a table t with match kind k , and h is not guaranteed to be valid at the call
655 site of t , and if $h \notin \text{valid_reads}(t)$ or the match-kind of $k \neq \text{ternary}$.

656 In the case of the `port_vlan_mapping` table in Figure 6, there is a valid bit for both
657 `vlan_tag_[0]` and `vlan_tag_[1]`, both of which are followed by `exact` matches. To solve
658 this problem, we need to use the `ternary` match-kind instead, which allows the use of
659 wildcard matching. When a field is matched with a wildcard, the table does not attempt to
660 compute the value of the `reads` expression; instead, the table short-circuits and skips the
661 check entirely.

662 This fix assumes that the controller is well behaved and fills the `vlan_tag_[0].vid` with
663 a wildcard whenever `vlan_tag_[0]` is matched as invalid (and similarly for `vlan_tag_[1]`)s.
664 This also what the `SAFE P4` type system does. `P4CHECK` prints warnings describing these
665 assumptions to the programmer (top of Figure 21), giving them properties against which to
666 check their control plane implementation.

667 4.2.4 Table Action Bugfixes

668 The table actions bugs, which are exemplified by the table `fabric_ingress_dst_lkp` from
669 Figure 7, can be fixed by modifying the `reads` declaration in the table. Recall that the
670 parser will parse exactly one of the headers `fabric_hdr_cpu`, `fabric_hdr_unicast` and
671 `fabric_hdr_multicast`, which means that when the table is applied at type Θ , exactly one
672 of `Includes Θ fabric_hdr_i` for $i \in \{\text{cpu}, \text{unicast}, \text{multicast}\}$ will hold. Now, the action
673 `term_cpu_packet` typechecks only with the (nonempty) type `Restrict Θ fabric_hdr_cpu`,
674 and the actions `term_fabric_i_packet` only typecheck with the (nonempty) types
675 `Restrict Θ term_fabric_i_packet` for $i = \text{unicast}, \text{multicast}$. `P4CHECK` suggests that
676 this is the cause of the bug since it reports type violations for all of the references to these three
677 headers in the control paths following from the application of `fabric_ingress_dst_lkp`.

678 The optimal⁵ fix here is to augment the `reads` declaration to include a validity check for

⁵ The other fix would be to refactor the single into multiple tables, each of which is guarded by separate validity checks. However, combining this kind of logic in a single table helps to conserve on-switch memory, so in striving to change the behavior of the program as little as possible, we propose modifying the table reads.

each contentious header. We then assume that the controller is well-behaved enough to only call actions when their required headers are valid, allowing us to typecheck each action in the appropriate type restriction. P4CHECK alerts the programmer whenever it makes such an assumption. We show these warnings for the fixed version of `fabric_ingress_dst_lkp` below the line in Figure 21.

4.2.5 Default Action Bugfix

The default action bug occurs when a programmer creates a wrapper table for an action that modifies the type, and forgets to force the table to call that action when the packet misses. The `add_value_header_1` table from Figure 8 wraps the action `add_value_header_1_act`, which calls the single line `add_header(nc_value_1)`.

The default action, when left unspecified, is `nop`, which means that if the pre-application type was Θ , then the post-application type is $\Theta + \Theta \cdot \text{nc_value_1}$, which does not include `nc_value_1`. Hence, P4CHECK reports every subsequent reference (on this code path) to `nc_header_1` to be a type violation.

To fix this bug, we need to set the default action to `add_value_1`—this makes the post-application type $\Theta \cdot \text{nc_value_1} + \Theta \cdot \text{nc_value_1} = \Theta \cdot \text{nc_value_1}$, which includes `nc_value_1`, thus allowing the subsequent code to typecheck.

4.3 Overhead

It is important to evaluate two kinds of overhead when considering a static type system: overhead on programmers and on the underlying implementation.

Typically, adding a static type system to a dynamic type system requires more work for the programmer—the field of gradual typing is devoted breaking the gargantuan task into smaller commit-sized chunks [5]. Surprisingly, in our experience, migrating real-world P4 code to pass the SAFEP4 type system required only modest programmer effort.

To qualitatively evaluate the effort required to change an unsafe program into a safe one using our type system, we manually fixed all of the bugs that we detected. The programs that had bugs required us to edit between 0.10% and 1.4% of the lines of code. The one exception was `PPPoE_USING_P4`, which was a 143 line program that required 6 line-edits (4%), all of which were validity checks. Conversely, `switch.p4` required 34 line edits, the greatest observed number, but this only accounted for 0.37% of the total lines of code in the program.

Each class of bugs has a simple one-to-two line fix, as described in Section 4.2: adding a validity check, adding a default action, or slightly from the parser. Each of these changes was straightforward to identify and simple to make.

Another possible concern is that that extending tables with extra read expressions, or adding run-time validity checks to controls, might impose a heavy cost on implementations, especially on hardware. Although we have not yet performed an extensive study of the impact on compiled code, based on the size and complexity of the annotations we added, we believe the additional cost should be quite low. Overall, given the large number of potential bugs located by P4CHECK, we believe the assurance one gains about safety properties by using a static type system makes the costs well worth it.

5 Related Work

Probably the most closely related work to SAFE4 is `p4v` [18]. Unlike SAFE4, which is based on a static type system, `p4v` uses Dijkstra’s approach to program verification based on predicate transformer semantics. To model the behavior of the control plane, `p4v` uses first-order annotations. SAFE4’s typing rule for table application is inspired by this idea, but adopts simple heuristics rather than requiring logical annotations.

Both tools be used to verify safety properties of data planes modelled in P4—e.g., that no read or write operations are possible on an invalid header. However, unlike SAFE4, `p4v` does not define a formal semantics of the P4 language and hence does not formally prove the soundness of the approach. As it is often the case when comparing approaches based on types to those based on program verification, `p4v` can check more complex properties, including architectural invariants and program-specific properties—e.g., that the IPv4 time-to-live field is correctly decremented on every packet. However, in general, it requires annotating the program with formal specifications both for the correctness property itself and to model the behavior of the control plane.

McKeown et al. developed an operational semantics for P4 [20], which is translated to Datalog to verify safety properties and to check program equivalence. An operational semantics for P4 was also developed in the K framework [25], yielding a symbolic model checker and deductive verification tool [14]. Vera [28] models the semantics of P4 by translation to SymNet [29], and develops a symbolic execution engine for verifying a variety of properties, including header validity.

Compared to SAFE4, these approaches do not use their formalization of P4 as a foundation for defining a type system that addresses common bugs. To the best of our knowledge, SAFE4 is the first formal calculus for a P4-like packet processing language that provides correct-by-construction guarantees of header safety properties.

Other languages have used type systems to rule out safety problems due to null references. For example, NullAway [27] analyzes all Java programs annotated with `@Nullable` annotations, making path-sensitive deductions about which references may be null. Similar to the validity checks in SAFE4, NullAway analyses conditionals for null checks of the form `var != null` using data flow analysis.

Looking further afield, PacLang [9] is a concurrent packet-processing language that uses a linear type system to allow multiple references to a given packet within a single thread. PacLang and SAFE4 share the use of a type system for verifying safety properties but they differ in the kind of properties they address and, hence, the kind of type system they employ for this purpose. In addition, the primary focus in PacLang is on efficient compilation whereas SAFE4 is concerned with ensuring the type safety of header data manipulated by the program.

Domino [26] is a domain-specific language for data plane algorithms supporting *packet transactions*—i.e., blocks of code that are guaranteed to be atomic and isolated from other transactions. In Domino, the programmer defines the operations needed for each packet without worrying about other in-flight packets. If it succeeds, the compiler guarantees performance at the line rate supported on programmable switches. Overall, Domino focuses on transactional guarantees and concurrency rather than header safety properties.

BPF+ [3] and eBPF [8] are packet-processing frameworks that can be used to extend the kernel networking stack with custom functionality. The modern eBPF framework is based on machine-level programming model, but it uses a virtual machine and code verifier to ensure a variety of basic safety properties. Much of the recent work on eBPF focuses on

techniques such as just-in-time compilation to achieve good performance.

SNAP [1] is a language for stateful packet processing based on P4. It offers a programming model with global state registers that are distributed across many physical switches while optimizing for various criteria, such as minimizing congestion. More specifically, the compiler analyses read/write dependencies to automatically optimize the placement of state and the routing of traffic across the underlying physical topology.

Of course, there is a long tradition of formal calculi that aim to capture some aspect of computation and make it amenable for mathematical reasoning. The design of SAFE4 is directly inspired by Featherweight Java [10], which stands out for its elegant formalization of a real-world languages in an extensible core calculus.

6 Conclusion

P4 provides a powerful programming model for network devices based on high-level and declarative abstractions. Unfortunately, P4 lacks basic safety guarantees, which often lead to a variety of bugs in practice. This paper proposes SAFE4, a domain-specific language for programmable data planes that comes equipped with a formal semantics and a static type system which ensures that every read or write to a header at run-time will be safe. Under the covers, SAFE4 uses a rich set of types that tracks dependencies between headers, as well as a path-sensitive analysis and domain-specific heuristics that model common idioms for programming control planes and minimize false positives. Our experiments using an OCaml prototype and a suite of open-source programs found on GitHub show that most P4 applications can be made safe with minimal programming effort. We hope that our work can help lay the foundation for future enhancements to P4 as well as the next generation of data plane languages. In the future, we plan to explore enriching SAFE4's type system to track additional properties, investigate correct-by-construction techniques for writing control-plane code, and develop a compiler for the language.

References

- 1 Mina Tahmasbi Arashloo, Yaron Koral, Michael Greenberg, Jennifer Rexford, and David Walker. Snap: Stateful network-wide abstractions for packet processing. In *Proceedings of the 2016 ACM SIGCOMM Conference*, SIGCOMM '16, pages 29–43, New York, NY, USA, 2016. ACM. URL: <http://doi.acm.org/10.1145/2934872.2934892>, doi:10.1145/2934872.2934892.
- 2 Jiasong Bai, Jun Bi, Menghao Zhang, and Guanyu Li. Filtering spoofed ip traffic using switching asics. In *Proceedings of the ACM SIGCOMM 2018 Conference on Posters and Demos*, pages 51–53. ACM, 2018.
- 3 Andrew Begel, Steven McCanne, and Susan L. Graham. Bpf+: Exploiting global data-flow optimization in a generalized packet filter architecture. In *Proceedings of the Conference on Applications, Technologies, Architectures, and Protocols for Computer Communication*, SIGCOMM '99, pages 123–134, New York, NY, USA, 1999. ACM. URL: <http://doi.acm.org/10.1145/316188.316214>, doi:10.1145/316188.316214.
- 4 Pat Bosshart, Dan Daly, Glen Gibb, Martin Izzard, Nick McKeown, Jennifer Rexford, Cole Schlesinger, Dan Talayco, Amin Vahdat, George Varghese, and David Walker. P4: Programming protocol-independent packet processors. *SIGCOMM Comput. Commun. Rev.*, 44(3):87–95, July 2014. URL: <http://doi.acm.org/10.1145/2656877.2656890>, doi:10.1145/2656877.2656890.
- 5 John Peter Campora, Sheng Chen, Martin Erwig, and Eric Walkingshaw. Migrating gradual types. *Proceedings of the ACM on Programming Languages*, 2(POPL):15, 2017.

- 813 **6** Martin Casado, Michael J Freedman, Justin Pettit, Jianying Luo, Nick McKeown, and
814 Scott Shenker. Ethane: Taking control of the enterprise. In *ACM SIGCOMM Computer*
815 *Communication Review*, volume 37, pages 1–12. ACM, 2007.
- 816 **7** P4 Language Consortium. P4 language specification, version 1.0.4. Technical report, Avail-
817 able at <https://p4.org/specs/>, 2017.
- 818 **8** Jonathan Corbet. Bpf: the universal in-kernel virtual machine, May 2014. Available at
819 <https://lwn.net/Articles/599755/>.
- 820 **9** Robert Ennals, Richard Sharp, and Alan Mycroft. Linear types for packet processing.
821 In David Schmidt, editor, *Programming Languages and Systems*, pages 204–218, Berlin,
822 Heidelberg, 2004. Springer Berlin Heidelberg.
- 823 **10** Atsushi Igarashi, Benjamin C. Pierce, and Philip Wadler. Featherweight java: A minimal
824 core calculus for java and gj. *ACM Trans. Program. Lang. Syst.*, 23(3):396–450, May 2001.
825 URL: <http://doi.acm.org/10.1145/503502.503505>, doi:10.1145/503502.503505.
- 826 **11** Xin Jin. netcache-p4, Mar 2018. URL: <https://github.com/netx-repo/netcache-p4>.
- 827 **12** Xin Jin, Xiaozhou Li, Haoyu Zhang, Nate Foster, Jeongkeun Lee, Robert Soulé, Changhoon
828 Kim, and Ion Stoica. NetChain: Scale-free sub-rtt coordination. In *USENIX Symposium*
829 *on Networked Systems Design and Implementation (NSDI)*, April 2018. Best paper award.
- 830 **13** Xin Jin, Xiaozhou Li, Haoyu Zhang, Robert Soulé, Jeongkeun Lee, Nate Foster, Changhoon
831 Kim, and Ion Stoica. Netcache: Balancing key-value stores with fast in-network caching. In
832 *Proceedings of the 26th Symposium on Operating Systems Principles*, pages 121–136. ACM,
833 2017.
- 834 **14** Ali Kheradmand and Grigore Roşu. P4k: A formal semantics of p4 and applications. Techni-
835 cal Report <https://arxiv.org/abs/1804.01468>, University of Illinois at Urbana-Champaign,
836 April 2018.
- 837 **15** George T. Klees, Andrew Ruef, Benjamin Cooper, Shiyi Wei, and Michael Hicks. Evaluating
838 fuzz testing. In *Proceedings of the ACM Conference on Computer and Communications*
839 *Security (CCS)*, October 2018.
- 840 **16** Chaitanya Kodeboyina. An open-source p4 switch with sai support, Jun 2015. URL:
841 <https://p4.org/p4/an-open-source-p4-switch-with-sai-support.html>.
- 842 **17** Rahul Kumar and BB Gupta. Stepping stone detection techniques: Classification and state-
843 of-the-art. In *Proceedings of the international conference on recent cognizance in wireless*
844 *communication & image processing*, pages 523–533. Springer, 2016.
- 845 **18** Jed Liu, William Hallahan, Cole Schlesinger, Milad Sharif, Jeongkeun Lee, Robert Soulé,
846 Han Wang, Călin Caşcaval, Nick McKeown, and Nate Foster. P4v: Practical verification
847 for programmable data planes. In *Proceedings of the 2018 Conference of the ACM Special*
848 *Interest Group on Data Communication*, SIGCOMM ’18, pages 490–503, New York, NY,
849 USA, 2018. ACM. URL: <http://doi.acm.org/10.1145/3230543.3230582>, doi:10.1145/
850 3230543.3230582.
- 851 **19** Nick McKeown, Tom Anderson, Hari Balakrishnan, Guru Parulkar, Larry Peterson, Jen-
852 nifer Rexford, Scott Shenker, and Jonathan Turner. Openflow: Enabling innovation in
853 campus networks. *SIGCOMM Comput. Commun. Rev.*, 38(2):69–74, March 2008. URL:
854 <http://doi.acm.org/10.1145/1355734.1355746>, doi:10.1145/1355734.1355746.
- 855 **20** Nick McKeown, Dan Talayco, George Varghese, Nuno Lopes, Niko-
856 laj Björner, and Andrey Rybalchenko. Automatically verifying reacha-
857 bility and well-formedness in p4 networks. Technical report, September
858 2016. URL: [https://www.microsoft.com/en-us/research/publication/](https://www.microsoft.com/en-us/research/publication/automatically-verifying-reachability-well-formedness-p4-networks/)
859 [automatically-verifying-reachability-well-formedness-p4-networks/](https://www.microsoft.com/en-us/research/publication/automatically-verifying-reachability-well-formedness-p4-networks/).
- 860 **21** Robin Milner. A theory of type polymorphism in programming. *Journal of Computer and*
861 *System Sciences*, 17(3):348–375, dec 1978.

- 22 Barefoot Networks. Behavioral model, Dec 2018. URL: <https://github.com/p4lang/behavioral-model>.
- 23 Barefoot Networks. Switch, Jan 2018. URL: <https://github.com/p4lang/switch>.
- 24 Tj OConnor, William Enck, W Michael Petullo, and Akash Verma. Pivotwall: Sdn-based information flow control. In *Proceedings of the Symposium on SDN Research*, page 3. ACM, 2018.
- 25 Grigore Roşu and Traian Florin Şerbănuţă. An overview of the K semantic framework. *Journal of Logic and Algebraic Programming*, 79(6):397–434, 2010. doi:10.1016/j.jlap.2010.03.012.
- 26 Anirudh Sivaraman, Alvin Cheung, Mihai Budiu, Changhoon Kim, Mohammad Alizadeh, Hari Balakrishnan, George Varghese, Nick McKeown, and Steve Licking. Packet transactions: High-level programming for line-rate switches. In *Proceedings of the 2016 ACM SIGCOMM Conference*, SIGCOMM '16, pages 15–28, New York, NY, USA, 2016. ACM. URL: <http://doi.acm.org/10.1145/2934872.2934900>, doi:10.1145/2934872.2934900.
- 27 Manu Sridharan. Engineering nullaway, uber’s open source tool for detecting nullpointerexceptions on android, Dec 2018. URL: <https://eng.uber.com/nullaway/>.
- 28 Radu Stoenescu, Dragos Dumitrescu, Matei Popovici, Lorina Negreanu, and Costin Raiciu. Debugging p4 programs with Vera. In *ACM SIGCOMM*, pages 518–532, New York, NY, USA, 2018. ACM. URL: <http://doi.acm.org/10.1145/3230543.3230548>, doi:10.1145/3230543.3230548.
- 29 Radu Stoenescu, Matei Popovici, Lorina Negreanu, and Costin Raiciu. Symnet: Scalable symbolic execution for modern networks. In *ACM SIGCOMM*, pages 314–327, New York, NY, USA, 2016. ACM. URL: <http://doi.acm.org/10.1145/2934872.2934881>, doi:10.1145/2934872.2934881.
- 30 Sam Tobin-Hochstadt and Matthias Felleisen. Logical types for untyped languages. In *Proceedings of the 15th ACM SIGPLAN International Conference on Functional Programming*, ICFP '10, pages 117–128, New York, NY, USA, 2010. ACM. URL: <http://doi.acm.org/10.1145/1863543.1863561>, doi:10.1145/1863543.1863561.
- 31 Menghao Zhang. Anti-spoof, Nov 2018. URL: <https://github.com/zhangmenghao/Anti-spoof>.

A Additional Operational Semantics rules

This section presents additional evaluation rules.

$$\begin{array}{c} \text{E-SEQ} \\ \hline \langle I, O, H, \text{skip}; c_2 \rangle \rightarrow \langle I, O, H, c_2 \rangle \end{array} \qquad \begin{array}{c} \text{E-SEQ1} \\ \hline \frac{\langle I, O, H, c_1 \rangle \rightarrow \langle I', O', H', c'_1 \rangle}{\langle I, O, H, c_1; c_2 \rangle \rightarrow \langle I', O', H', c'_1; c_2 \rangle} \end{array}$$

$$\begin{array}{c} \text{E-IF} \\ \hline \frac{\langle H, e \rangle \rightarrow e'}{\langle I, O, H, \text{if } (e) \text{ then } c_1 \text{ else } c_2 \rangle \rightarrow \langle I, O, H, \text{if } (e') \text{ then } c_1 \text{ else } c_2 \rangle} \end{array}$$

E-IFTRUE

$$\hline \langle I, O, H, \text{if } (true) \text{ then } c_1 \text{ else } c_2 \rangle \rightarrow \langle I, O, H, c_1 \rangle$$

E-IFFALSE

$$\hline \langle I, O, H, \text{if } (false) \text{ then } c_1 \text{ else } c_2 \rangle \rightarrow \langle I, O, H, c_2 \rangle$$

893 **B** Operations on header types

894 This section presents an in-depth treatment of the operations defined on header instance
895 types. In the following we assume that S ranges over elements of the domain $\mathcal{P}(\mathcal{P}(H))$.

896 **Restriction** The restrict operator $\text{Restrict } \Theta \ h$ recursively traverses Θ and keeps only those
897 choices in which h is contained, zeroing out the others. Semantically this has the effect of
898 throwing out the subsets of $\llbracket \Theta \rrbracket$ that do not contain h , i.e., we define restriction semantically
899 as $S|h \triangleq \{hs | hs \in S \wedge h \in hs\}$. Syntactically we define restriction by induction on Θ as
shown in Figure 22. The equivalence of the syntactic and the semantic definition is captured

$$\begin{aligned}
 \text{Restrict } 0 \ h &\triangleq 0 \\
 \text{Restrict } 1 \ h &\triangleq 0 \\
 \text{Restrict } g \ h &\triangleq \begin{cases} g & \text{if } g = h \\ 0 & \text{otherwise} \end{cases} \\
 \text{Restrict } (\Theta_1 \cdot \Theta_2) \ h &\triangleq ((\text{Restrict } \Theta_1 \ h) \cdot \Theta_2) + (\Theta_1 \cdot (\text{Restrict } \Theta_2 \ h)) \\
 \text{Restrict } (\Theta_1 + \Theta_2) \ h &\triangleq (\text{Restrict } \Theta_1 \ h) + (\text{Restrict } \Theta_2 \ h)
 \end{aligned}$$

■ **Figure 22** Syntactic definition of the $\text{Restrict } \Theta \ h$ operator.

900
901 by Lemma 3.

902 ► **Lemma 3.** $\llbracket \Theta \rrbracket|h == \llbracket \text{Restrict } \Theta \ h \rrbracket$.

903 **Proof.** By induction on Θ .

904 *Case* $\Theta = 0$:

$$\begin{aligned}
 905 \quad &\llbracket 0 \rrbracket|h \\
 906 \quad &= \{\} | h && \text{by definition of } \llbracket \cdot \rrbracket \\
 907 \quad &= \{hs | hs \in \{\} \wedge h \in hs\} && \text{by definition of } \cdot | h \\
 908 \quad &= \{\} && \text{by set theory} \\
 909 \quad &= \llbracket 0 \rrbracket && \text{by definition of } \llbracket \cdot \rrbracket \\
 910 \quad &= \llbracket \text{Restrict } 0 \ h \rrbracket && \text{by definition of } \text{Restrict } \cdot h \\
 911
 \end{aligned}$$

912 *Case* $\Theta = 1$:

$$\begin{aligned}
 913 \quad &\llbracket 1 \rrbracket|h \\
 914 \quad &= \{\{\} \} | h && \text{by definition of } \llbracket \cdot \rrbracket \\
 915 \quad &= \{hs | hs \in \{\{\} \} \wedge h \in hs\} && \text{by definition of } \cdot | h \\
 916 \quad &= \{\} && \text{by set theory} \\
 917 \quad &= \llbracket 0 \rrbracket && \text{by definition of } \llbracket \cdot \rrbracket \\
 918 \quad &= \llbracket \text{Restrict } 1 \ h \rrbracket && \text{by definition of } \text{Restrict } \cdot h \\
 919
 \end{aligned}$$

920 *Case* $\Theta = g$:

$$\begin{aligned}
921 & \llbracket g \rrbracket | h \\
922 & = \{\{g\}\} | h && \text{by definition of } \llbracket \cdot \rrbracket \\
923 & = \{hs | hs \in \{\{g\}\} \wedge h \in hs\} && \text{by definition of } \cdot | h \\
924 & \quad \textit{Subcase } h = g \\
925 & = \{\{g\}\} && \text{by set theory} \\
926 & = \llbracket g \rrbracket && \text{by definition of } \llbracket \cdot \rrbracket \\
927 & = \llbracket \text{Restrict } g \ h \rrbracket && \text{by assumption } h = g \text{ and by definition of } \text{Restrict } \cdot h \\
928 & \quad \textit{Subcase } h \neq g \\
929 & = \{\} && \text{by set theory} \\
930 & = \llbracket 0 \rrbracket && \text{by definition of } \llbracket \cdot \rrbracket \\
931 & = \llbracket \text{Restrict } g \ h \rrbracket && \text{by definition of } \text{Restrict } \cdot h \text{ and by assumption } h \neq g
\end{aligned}$$

932
933

934 *Case* $\Theta = \Theta_1 \cdot \Theta_2$:

935

$$\begin{aligned}
936 & \llbracket \Theta_1 \cdot \Theta_2 \rrbracket | h \\
937 & = \{hs_1 \cup hs_2 | hs_1 \in \llbracket \Theta_1 \rrbracket \wedge hs_2 \in \llbracket \Theta_2 \rrbracket\} | h && \text{by definition of } \llbracket \cdot \rrbracket \\
938 & = \{hs_1 \cup hs_2 | hs_1 \in \llbracket \Theta_1 \rrbracket \wedge hs_2 \in \llbracket \Theta_2 \rrbracket \wedge h \in (hs_1 \cup hs_2)\} && \text{by definition of } \cdot | h \\
939 & = \{hs_1 \cup hs_2 | hs_1 \in \llbracket \Theta_1 \rrbracket \wedge hs_2 \in \llbracket \Theta_2 \rrbracket \wedge h \in hs_1\} \cup && \text{by set theory} \\
940 & \quad \{hs_1 \cup hs_2 | hs_1 \in \llbracket \Theta_1 \rrbracket \wedge hs_2 \in \llbracket \Theta_2 \rrbracket \wedge h \in hs_2\} \\
941 & = \{hs_1 \cup hs_2 | (hs_1 \in \llbracket \Theta_1 \rrbracket \wedge h \in hs_1) \wedge hs_2 \in \llbracket \Theta_2 \rrbracket\} \cup && \text{by logic and set theory} \\
942 & \quad \{hs_1 \cup hs_2 | hs_1 \in \llbracket \Theta_1 \rrbracket \wedge (hs_2 \in \llbracket \Theta_2 \rrbracket \wedge h \in hs_2)\} \\
943 & = \{hs_1 | hs_1 \in \llbracket \Theta_1 \rrbracket \wedge h \in hs_1\} \bullet \{hs_2 | hs_2 \in \llbracket \Theta_2 \rrbracket\} \cup && \text{by definition of } S_1 \bullet S_2 \\
944 & \quad \{hs_1 | hs_1 \in \llbracket \Theta_1 \rrbracket\} \bullet \{hs_2 | hs_2 \in \llbracket \Theta_2 \rrbracket \wedge h \in hs_2\} \\
945 & = \{hs_1 | hs_1 \in \llbracket \Theta_1 \rrbracket\} | h \bullet \{hs_2 | hs_2 \in \llbracket \Theta_2 \rrbracket\} \cup && \text{by definition of } \cdot | h \\
946 & \quad \{hs_1 | hs_1 \in \llbracket \Theta_1 \rrbracket\} \bullet \{hs_2 | hs_2 \in \llbracket \Theta_2 \rrbracket\} | h \\
947 & = \llbracket \Theta_1 \rrbracket | h \bullet \llbracket \Theta_2 \rrbracket \cup \llbracket \Theta_1 \rrbracket \bullet \llbracket \Theta_2 \rrbracket | h && \text{by definition of } \llbracket \cdot \rrbracket \\
948 & = \llbracket \text{Restrict } \Theta_1 \ h \rrbracket \bullet \llbracket \Theta_2 \rrbracket \cup \llbracket \Theta_1 \rrbracket \bullet \llbracket \text{Restrict } \Theta_2 \ h \rrbracket && \text{by induction hypothesis} \\
949 & = \llbracket \text{Restrict } \Theta_1 \ h \cdot \Theta_2 + \Theta_1 \cdot \text{Restrict } \Theta_2 \ h \rrbracket && \text{by definition of } S_1 \bullet S_2 \text{ and } \llbracket \cdot \rrbracket \\
950 & = \llbracket \text{Restrict } (\Theta_1 \cdot \Theta_2) \ h \rrbracket && \text{by definition of } \text{Restrict } \cdot h
\end{aligned}$$

952 *Case* $\Theta = \Theta_1 + \Theta_2$:

$$\begin{aligned}
953 & \llbracket \Theta_1 + \Theta_2 \rrbracket | h \\
954 & = (\llbracket \Theta_1 \rrbracket \cup \llbracket \Theta_2 \rrbracket) | h && \text{by definition of } \llbracket \cdot \rrbracket \\
955 & = \{hs | hs \in (\llbracket \Theta_1 \rrbracket \cup \llbracket \Theta_2 \rrbracket) \wedge h \in hs\} && \text{by definition of } \cdot | h \\
956 & = \{hs_1 | hs_1 \in \llbracket \Theta_1 \rrbracket \wedge h \in hs_1\} \cup \{hs_2 | hs_2 \in \llbracket \Theta_2 \rrbracket \wedge h \in hs_2\} && \text{by set theory} \\
957 & = \llbracket \Theta_1 \rrbracket | h \cup \llbracket \Theta_2 \rrbracket | h && \text{by definition of } \cdot | h \\
958 & = \llbracket \text{Restrict } \Theta_1 \ h \rrbracket \cup \llbracket \text{Restrict } \Theta_2 \ h \rrbracket && \text{by induction hypothesis} \\
959 & = \llbracket \text{Restrict } \Theta_1 \ h + \text{Restrict } \Theta_2 \ h \rrbracket && \text{by definition of } \llbracket \cdot \rrbracket \\
960 & = \llbracket \text{Restrict } (\Theta_1 + \Theta_2) \ h \rrbracket && \text{by definition of } \text{Restrict } \cdot h
\end{aligned}$$

962



963 **Negated Restriction** Dually to the restrict operator, `NegRestrict` Θ h produces only
 964 those choices/subsets where h is invalid. Semantically, negated restriction is defined as
 965 $S|\neg h \triangleq \{hs | hs \in S \wedge h \notin hs\}$. Syntactically we define *Negated Restriction* by induction on
 Θ as shown in Figure 23.

$$\begin{aligned}
 \text{NegRestrict } 0 \ h &\triangleq 0 \\
 \text{NegRestrict } 1 \ h &\triangleq 1 \\
 \text{NegRestrict } g \ h &\triangleq \begin{cases} 0 & \text{if } g = h \\ g & \text{otherwise} \end{cases} \\
 \text{NegRestrict } (\Theta_1 \cdot \Theta_2) \ h &\triangleq (\text{NegRestrict } \Theta_1 \ h) \cdot (\text{NegRestrict } \Theta_2 \ h) \\
 \text{NegRestrict } (\Theta_1 + \Theta_2) \ h &\triangleq (\text{NegRestrict } \Theta_1 \ h) + (\text{NegRestrict } \Theta_2 \ h)
 \end{aligned}$$

■ **Figure 23** Syntactic definition of the `NegRestrict` Θ h operator

966

967 The equivalence of the syntactic and semantic definition is captured by Lemma 4.

968 ► **Lemma 4.** $\llbracket \Theta \rrbracket |\neg h == \llbracket \text{NegRestrict } \Theta \ h \rrbracket$.

969 **Proof.** By induction on Θ .

970 *Case* $\Theta = 0$:

$$\begin{aligned}
 971 \quad &\llbracket 0 \rrbracket |\neg h \\
 972 \quad &= \{\} |\neg h && \text{by definition of } \llbracket . \rrbracket \\
 973 \quad &= \{hs | hs \in \{\} \wedge h \notin hs\} && \text{by definition of } .|\neg h \\
 974 \quad &= \{\} && \text{by set theory} \\
 975 \quad &= \llbracket 0 \rrbracket && \text{by definition of } \llbracket . \rrbracket \\
 976 \quad &= \llbracket \text{NegRestrict } 0 \ h \rrbracket && \text{by definition of } \text{NegRestrict} . h \\
 977
 \end{aligned}$$

978 *Case* $\Theta = 1$:

$$\begin{aligned}
 979 \quad &\llbracket 1 \rrbracket |\neg h \\
 980 \quad &= \{\{\}\} |\neg h && \text{by definition of } \llbracket . \rrbracket \\
 981 \quad &= \{hs | hs \in \{\{\}\} \wedge h \notin hs\} && \text{by definition of } .|\neg h \\
 982 \quad &= \{\{\}\} && \text{by set theory} \\
 983 \quad &= \llbracket 1 \rrbracket && \text{by definition of } \llbracket . \rrbracket \\
 984 \quad &= \llbracket \text{NegRestrict } 1 \ h \rrbracket && \text{by definition of } \text{NegRestrict} . h \\
 985
 \end{aligned}$$

986 *Case* $\Theta = g$:

$$\begin{aligned}
987 & \llbracket g \rrbracket | \neg h \\
988 & = \{\{g\}\} | \neg h && \text{by definition of } \llbracket \cdot \rrbracket \\
989 & = \{hs | hs \in \{\{g\}\} \wedge h \notin hs\} && \text{by definition of } \cdot | \neg h \\
990 & \quad \textit{Subcase } h = g \\
991 & = \{\} && \text{by set theory} \\
992 & = \llbracket 0 \rrbracket && \text{by definition of } \llbracket \cdot \rrbracket \\
993 & = \llbracket \text{NegRestrict } 0 \ h \rrbracket && \text{by assumption } h = g \text{ and by definition of } \text{NegRestrict} \cdot h \\
994 & \quad \textit{Subcase } h \neq g \\
995 & = \{\{g\}\} && \text{by set theory} \\
996 & = \llbracket g \rrbracket && \text{by definition of } \llbracket \cdot \rrbracket \\
997 & = \llbracket \text{NegRestrict } g \ h \rrbracket && \text{by definition of } \text{NegRestrict} \cdot h \text{ and by assumption } h \neq g
\end{aligned}$$

998

1000 *Case* $\Theta = \Theta_1 \cdot \Theta_2$:

1001

$$\begin{aligned}
1002 & \llbracket \Theta_1 \cdot \Theta_2 \rrbracket | \neg h \\
1003 & = (\llbracket \Theta_1 \rrbracket \bullet \llbracket \Theta_2 \rrbracket) | \neg h && \text{by definition of } S_1 \bullet S_2 \\
1004 & = \{hs_1 \cup hs_2 | hs_1 \in \llbracket \Theta_1 \rrbracket \wedge hs_2 \in \llbracket \Theta_2 \rrbracket\} | \neg h && \text{by definition of } \llbracket \cdot \rrbracket \\
1005 & = \{hs_1 \cup hs_2 | hs_1 \in \llbracket \Theta_1 \rrbracket \wedge hs_2 \in \llbracket \Theta_2 \rrbracket \wedge h \notin (hs_1 \cup hs_2)\} && \text{by definition of } \cdot | \neg h \\
1006 & = \{hs_1 \cup hs_2 | hs_1 \in \llbracket \Theta_1 \rrbracket \wedge hs_2 \in \llbracket \Theta_2 \rrbracket \wedge h \notin hs_1 \wedge h \notin hs_2\} && \text{by set theory and logic} \\
1007 & = \{hs_1 \cup hs_2 | (hs_1 \in \llbracket \Theta_1 \rrbracket \wedge h \notin hs_1) \wedge (hs_2 \in \llbracket \Theta_2 \rrbracket \wedge h \notin hs_2)\} && \text{by set theory and logic} \\
1008 & = \{hs_1 | hs_1 \in \llbracket \Theta_1 \rrbracket \wedge h \notin hs_1\} \bullet \{hs_2 | hs_2 \in \llbracket \Theta_2 \rrbracket \wedge h \notin hs_2\} && \text{by definition of } S_1 \bullet S_2 \\
1009 & = \llbracket \Theta_1 \rrbracket | \neg h \bullet \llbracket \Theta_2 \rrbracket | \neg h && \text{by definition of } \cdot | \neg h \\
1010 & = \llbracket \text{NegRestrict } \Theta_1 \ h \rrbracket \bullet \llbracket \text{NegRestrict } \Theta_2 \ h \rrbracket && \text{by induction hypothesis} \\
1011 & = \llbracket (\text{NegRestrict } \Theta_1 \ h) \cdot (\text{NegRestrict } \Theta_2 \ h) \rrbracket && \text{By definition of } \llbracket \cdot \rrbracket \\
1012 & = \llbracket \text{NegRestrict } (\Theta_1 \cdot \Theta_2) \ h \rrbracket && \text{by definition of } \text{NegRestrict} \cdot h
\end{aligned}$$

1013

1015 *Case* $\Theta = \Theta_1 + \Theta_2$:

$$\begin{aligned}
1016 & \llbracket \Theta_1 + \Theta_2 \rrbracket | \neg h \\
1017 & = (\llbracket \Theta_1 \rrbracket \cup \llbracket \Theta_2 \rrbracket) | \neg h && \text{by definition of } \llbracket \cdot \rrbracket \\
1018 & = \{hs | hs \in (\llbracket \Theta_1 \rrbracket \cup \llbracket \Theta_2 \rrbracket) \wedge h \notin hs\} && \text{by definition of } \cdot | \neg h \\
1019 & = \{hs_1 | hs_1 \in \llbracket \Theta_1 \rrbracket \wedge h \notin hs_1\} \cup \{hs_2 | hs_2 \in \llbracket \Theta_2 \rrbracket \wedge h \notin hs_2\} && \text{by set theory} \\
1020 & = \llbracket \Theta_1 \rrbracket | \neg h \cup \llbracket \Theta_2 \rrbracket | \neg h && \text{by definition of } \cdot | \neg h \\
1021 & = \llbracket \text{NegRestrict } \Theta_1 \ h \rrbracket \cup \llbracket \text{NegRestrict } \Theta_2 \ h \rrbracket && \text{by induction hypothesis} \\
1022 & = \llbracket \text{NegRestrict } \Theta_1 \ h + \text{NegRestrict } \Theta_2 \ h \rrbracket && \text{by definition of } \llbracket \cdot \rrbracket \\
1023 & = \llbracket \text{NegRestrict } (\Theta_1 + \Theta_2) \ h \rrbracket && \text{by definition of } \text{NegRestrict} \cdot \neg h
\end{aligned}$$

1025



1026 **Inclusion** Inclusion $\Theta \ h$ traverses Θ and checks to make sure that h is valid in every path.
 1027 Semantically this says that h is a member of every element of $\llbracket \Theta \rrbracket$, i.e., $h \sqsubset S \triangleq \bigwedge (hs \in S \wedge h \in hs)$. Syntactically we define *Inclusion* by induction on Θ as shown in Figure 24.

$$\begin{aligned}
 \text{Includes } 0 \ h &\triangleq \text{false} \\
 \text{Includes } 1 \ h &\triangleq \text{false} \\
 \text{Includes } g \ h &\triangleq \begin{cases} \text{true} & \text{if } g = h \\ \text{false} & \text{otherwise} \end{cases} \\
 \text{Includes } (\Theta_1 \cdot \Theta_2) \ h &\triangleq (\text{Includes } \Theta_1 \ h) \vee (\text{Includes } \Theta_2 \ h) \\
 \text{Includes } (\Theta_1 + \Theta_2) \ h &\triangleq (\text{Includes } \Theta_1 \ h) \wedge (\text{Includes } \Theta_2 \ h)
 \end{aligned}$$

■ **Figure 24** Syntactic definition of the **Includes** $\Theta \ h$ operator

1028 The equivalence of the syntactic and semantic definition is captured by Lemma 5.
 1029

1030 ► **Lemma 5.** $\forall hs \in S. h \in hs == \text{Includes } \Theta \ h.$

1031 **Proof.** By induction on Θ .

1032 *Case* $\Theta = 0$:

$$\begin{aligned}
 1033 \quad & h \sqsubset \llbracket 0 \rrbracket \\
 1034 \quad & = \bigwedge (hs \in \{\} \wedge h \in hs) && \text{by definition of } \llbracket \cdot \rrbracket \\
 1035 \quad & = \text{false} && \text{by logic and set theory} \\
 1036 \quad & = \text{Includes } 0 \ h && \text{by definition of Includes } \cdot \ h
 \end{aligned}$$

1038 *Case* $\Theta = 1$:

$$\begin{aligned}
 1039 \quad & h \sqsubset \llbracket 1 \rrbracket \\
 1040 \quad & = \bigwedge (hs \in \{\{\}\} \wedge h \in hs) && \text{by definition of } \llbracket \cdot \rrbracket \\
 1041 \quad & = \text{false} && \text{by logic and set theory} \\
 1042 \quad & = \text{Includes } 1 \ h && \text{by definition of Includes } \cdot \ h
 \end{aligned}$$

1044 *Case* $\Theta = g$:

$$\begin{aligned}
 1045 \quad & h \sqsubset \llbracket g \rrbracket \\
 1046 \quad & = \bigwedge (hs \in \{\{g\}\} \wedge h \in hs) && \text{by definition of } \llbracket \cdot \rrbracket \\
 1047 \quad & \text{Subcase } h = g \\
 1048 \quad & = \text{true} && \text{by logic and set theory} \\
 1049 \quad & = \text{Includes } g \ h && \text{by definition of (Includes } \cdot \ h) \text{ and assumption } h = g \\
 1050 \quad & \text{Subcase } h \neq g \\
 1051 \quad & = \text{false} && \text{by logic and set theory} \\
 1052 \quad & = \text{Includes } g \ h && \text{by definition of (Includes } \cdot \ h) \text{ and assumption } h \neq g
 \end{aligned}$$

1053
1054

1055 *Case* $\Theta = \Theta_1 \cdot \Theta_2$:

$$\begin{aligned}
1056 \quad & h \sqsubset \llbracket \Theta_1 \cdot \Theta_2 \rrbracket \\
1057 \quad & = h \sqsubset (\llbracket \Theta_1 \rrbracket \bullet \llbracket \Theta_2 \rrbracket) && \text{by definition of } S_1 \bullet S_2 \\
1058 \quad & = h \sqsubset \{hs_1 \cup hs_2 \mid hs_1 \in \llbracket \Theta_1 \rrbracket \wedge hs_2 \in \llbracket \Theta_2 \rrbracket\} && \text{by definition of } \llbracket \cdot \rrbracket \\
1059 \quad & = h \sqsubset \{hs_1 \mid hs_1 \in \llbracket \Theta_1 \rrbracket\} \vee h \sqsubset \{hs_2 \mid hs_2 \in \llbracket \Theta_2 \rrbracket\} && \text{by set theory and logic} \\
1060 \quad & = \bigwedge (hs_1 \in \llbracket \Theta_1 \rrbracket \wedge h \in hs_1) \vee \bigwedge (hs_2 \in \llbracket \Theta_2 \rrbracket \wedge h \in hs_2) && \text{by definition of } h \sqsubset \cdot \\
1061 \quad & = h \sqsubset \llbracket \Theta_1 \rrbracket \vee h \sqsubset \llbracket \Theta_2 \rrbracket && \text{by definition of } \llbracket \cdot \rrbracket \\
1062 \quad & = (\text{Includes } \Theta_1 \ h) \vee (\text{Includes } \Theta_2 \ h) && \text{by induction hypothesis} \\
1063 \quad & = (\text{Includes } \Theta_1 \cdot \Theta_2 \ h) && \text{by definition of Includes } \cdot \ h
\end{aligned}$$

1065 *Case* $\Theta = \Theta_1 + \Theta_2$:

$$\begin{aligned}
1066 \quad & h \sqsubset \llbracket \Theta_1 + \Theta_2 \rrbracket \\
1067 \quad & = h \sqsubset (\llbracket \Theta_1 \rrbracket \cup \llbracket \Theta_2 \rrbracket) && \text{by definition of } \llbracket \cdot \rrbracket \\
1068 \quad & = \bigwedge (hs \in (\llbracket \Theta_1 \rrbracket \cup \llbracket \Theta_2 \rrbracket) \wedge h \in hs) && \text{by definition of } h \sqsubset \cdot \\
1069 \quad & = \bigwedge (hs_1 \in \llbracket \Theta_1 \rrbracket \wedge h \in hs_1 \wedge hs_2 \in \llbracket \Theta_2 \rrbracket \wedge h \in hs_2) && \text{by set theory and logic} \\
1070 \quad & = \bigwedge (hs_1 \in \llbracket \Theta_1 \rrbracket \wedge h \in hs_1) \wedge \bigwedge (hs_2 \in \llbracket \Theta_2 \rrbracket \wedge h \in hs_2) && \text{by set theory and logic} \\
1071 \quad & = h \sqsubset \llbracket \Theta_1 \rrbracket \wedge h \sqsubset \llbracket \Theta_2 \rrbracket && \text{by definition of } h \sqsubset \cdot \\
1072 \quad & = (\text{Includes } \Theta_1 \ h) \wedge (\text{Includes } \Theta_2 \ h) && \text{by induction hypothesis} \\
1073 \quad & = (\text{Includes } (\Theta_1 + \Theta_2) \ h) && \text{by definition of Includes } \cdot \ h
\end{aligned}$$

1075

1076 **Removal** Remove $\Theta \ h$ removes h from every path, which means, semantically that it removes
 1077 h from every element of $\llbracket \Theta \rrbracket$, i.e., $S \setminus h \triangleq \{hs \mid hs \in S \wedge hs \setminus \{h\}\}$. Syntactically we define
Removal by induction on Θ as shown in Figure 25.

$$\begin{aligned}
& \text{Remove } 0 \ h \triangleq 0 \\
& \text{Remove } 1 \ h \triangleq 1 \\
& \text{Remove } g \ h \triangleq \begin{cases} 1 & \text{if } g = h \\ g & \text{otherwise} \end{cases} \\
& \text{Remove } (\Theta_1 \cdot \Theta_2) \ h \triangleq (\text{Remove } \Theta_1 \ h) \cdot (\text{Remove } \Theta_2 \ h) \\
& \text{Remove } (\Theta_1 + \Theta_2) \ h \triangleq (\text{Remove } \Theta_1 \ h) + (\text{Remove } \Theta_2 \ h)
\end{aligned}$$

■ **Figure 25** Syntactic definition of the Remove $\Theta \ h$ operator

1078

1079 The equivalence of the syntactic and semantic definition is captured by Lemma 6.

1080 ► **Lemma 6.** $\llbracket \Theta \rrbracket \setminus h == \llbracket \text{Remove } \Theta \ h \rrbracket$.

23:34 Type-Safe Data Plane Programming with SafeP4

1081 **Proof.** By induction on Θ .

1082 *Case* $\Theta = 0$:

$$\begin{aligned}
 1083 \quad & \llbracket 0 \rrbracket \setminus h \\
 1084 \quad &= \{\} \setminus h && \text{by definition of } \llbracket \cdot \rrbracket \\
 1085 \quad &= \{hs \mid hs \in \{\} \wedge hs \setminus \{h\}\} && \text{by definition of } \cdot \setminus h \\
 1086 \quad &= \{\} && \text{by set theory} \\
 1087 \quad &= \llbracket 0 \rrbracket && \text{by definition of } \llbracket \cdot \rrbracket \\
 1088 \quad &= \llbracket \text{Remove } 0 \ h \rrbracket && \text{by definition of } \text{Remove} \cdot h \\
 1089
 \end{aligned}$$

1090 *Case* $\Theta = 1$:

$$\begin{aligned}
 1091 \quad & \llbracket 1 \rrbracket \setminus h \\
 1092 \quad &= \{\{\}\} \setminus h && \text{by definition of } \llbracket \cdot \rrbracket \\
 1093 \quad &= \{hs \mid hs \in \{\{\}\} \wedge hs \setminus \{h\}\} && \text{by definition of } \cdot \setminus h \\
 1094 \quad &= \{\{\}\} && \text{by set theory} \\
 1095 \quad &= \llbracket 1 \rrbracket && \text{by definition of } \llbracket \cdot \rrbracket \\
 1096 \quad &= \llbracket \text{Remove } 1 \ h \rrbracket && \text{by definition of } \text{Remove} \cdot h \\
 1097
 \end{aligned}$$

1098 *Case* $\Theta = g$:

$$\begin{aligned}
 1099 \quad & \llbracket g \rrbracket \setminus h \\
 1100 \quad &= \{\{g\}\} \setminus h && \text{by definition of } \llbracket \cdot \rrbracket \\
 1101 \quad &= \{hs \mid hs \in \{\{g\}\} \wedge hs \setminus \{h\}\} && \text{by definition of } \cdot \setminus h \\
 1102 \quad & \quad \textit{Subcase } h = g \\
 1103 \quad &= \{\{\}\} && \text{by set theory} \\
 1104 \quad &= \llbracket 1 \rrbracket && \text{by definition of } \llbracket \cdot \rrbracket \\
 1105 \quad &= \llbracket \text{Remove } 1 \ h \rrbracket && \text{by definition of } \text{Remove} \cdot h \\
 1106 \quad & \quad \textit{Subcase } h \neq g \\
 1107 \quad &= \{\{g\}\} && \text{by set theory} \\
 1108 \quad &= \llbracket g \rrbracket && \text{by definition of } \llbracket \cdot \rrbracket \\
 1109 \quad &= \llbracket \text{Remove } g \ h \rrbracket && \text{by definition of } \text{Remove} \cdot h \\
 1110
 \end{aligned}$$

1111 *Case* $\Theta = \Theta_1 \cdot \Theta_2$:

$$\begin{aligned}
1112 & \llbracket \Theta_1 \cdot \Theta_2 \rrbracket \setminus h \\
1113 & = (\llbracket \Theta_1 \rrbracket \bullet \llbracket \Theta_2 \rrbracket) \setminus h && \text{by definition of } S_1 \bullet S_2 \\
1114 & = \{hs_1 \cup hs_2 \mid hs_1 \in \llbracket \Theta_1 \rrbracket \wedge hs_2 \in \llbracket \Theta_2 \rrbracket\} \setminus h && \text{by definition of } \llbracket \cdot \rrbracket \\
1115 & = \{hs_1 \cup hs_2 \mid hs_1 \in \llbracket \Theta_1 \rrbracket \wedge hs_2 \in \llbracket \Theta_2 \rrbracket \wedge (hs_1 \cup hs_2) \setminus h\} && \text{by definition of } \cdot \setminus h \\
1116 & = \{hs_1 \cup hs_2 \mid hs_1 \in \llbracket \Theta_1 \rrbracket \wedge hs_2 \in \llbracket \Theta_2 \rrbracket \wedge hs_1 \setminus h \wedge hs_2 \setminus h\} && \text{by set theory and logic} \\
1117 & = \{hs_1 \cup hs_2 \mid (hs_1 \in \llbracket \Theta_1 \rrbracket \wedge hs_1 \setminus h) \wedge (hs_2 \in \llbracket \Theta_2 \rrbracket \wedge hs_2 \setminus h)\} && \text{by set theory and logic} \\
1118 & = \{hs_1 \mid hs_1 \in \llbracket \Theta_1 \rrbracket \wedge hs_1 \setminus h\} \bullet \{hs_2 \mid hs_2 \in \llbracket \Theta_2 \rrbracket \wedge hs_2 \setminus h\} && \text{by definition of } S_1 \bullet S_2 \\
1119 & = \llbracket \Theta_1 \rrbracket \setminus h \bullet \llbracket \Theta_2 \rrbracket \setminus h && \text{by definition of } \cdot \setminus h \\
1120 & = \llbracket \text{Remove } \Theta_1 \ h \rrbracket \bullet \llbracket \text{Remove } \Theta_2 \ h \rrbracket && \text{by induction hypothesis} \\
1121 & = \llbracket (\text{Remove } \Theta_1 \ h) \cdot (\text{Remove } \Theta_2 \ h) \rrbracket && \text{By definition of } \llbracket \cdot \rrbracket \\
1122 & = \llbracket \text{Remove } (\Theta_1 \cdot \Theta_2) \ h \rrbracket && \text{by definition of } \text{Remove } \cdot \ h
\end{aligned}$$

1123

1125 *Case* $\Theta = \Theta_1 + \Theta_2$:

$$\begin{aligned}
1126 & \llbracket \Theta_1 + \Theta_2 \rrbracket \setminus h \\
1127 & = (\llbracket \Theta_1 \rrbracket \cup \llbracket \Theta_2 \rrbracket) \setminus h && \text{by definition of } \llbracket \cdot \rrbracket \\
1128 & = \{hs \mid hs \in (\llbracket \Theta_1 \rrbracket \cup \llbracket \Theta_2 \rrbracket) \wedge hs \setminus \{h\}\} && \text{by definition of } \cdot \setminus h \\
1129 & = \{hs \mid hs \in (\llbracket \Theta_1 \rrbracket \cup \llbracket \Theta_2 \rrbracket) \wedge hs \setminus \{h\}\} && \text{by definition of } \cdot \setminus h \\
1130 & = \{hs_1 \mid hs_1 \in \llbracket \Theta_1 \rrbracket \wedge hs_1 \setminus \{h\}\} \cup && \text{by logic and set theory} \\
1131 & \quad \{hs_2 \mid hs_2 \in \llbracket \Theta_2 \rrbracket \wedge hs_2 \setminus \{h\}\} \\
1132 & = \llbracket \Theta_1 \rrbracket \setminus h \cup \llbracket \Theta_2 \rrbracket \setminus h && \text{by definition of } \cdot \setminus h \\
1133 & = \llbracket \text{Remove } \Theta_1 \ h \rrbracket \cup \llbracket \text{Remove } \Theta_2 \ h \rrbracket && \text{by induction hypothesis} \\
1134 & = \llbracket (\text{Remove } \Theta_1 \ h) \cdot (\text{Remove } \Theta_2 \ h) \rrbracket && \text{by definition of } \llbracket \cdot \rrbracket \\
1135 & = \llbracket \text{Remove } (\Theta_1 + \Theta_2) \ h \rrbracket && \text{by definition of } \text{Remove } + \ h
\end{aligned}$$

1136

1137

1138

1139 **Emptiness** *Empty* Θ checks if Θ is semantically empty. Syntactically we define *Empty* by
 1140 induction on Θ as shown in Figure 26.

$ \begin{aligned} & \text{Empty } 0 \triangleq \text{true} \\ & \text{Empty } 1 \triangleq \text{false} \\ & \text{Empty } h \triangleq \text{false} \\ & \text{Empty } (\Theta_1 \cdot \Theta_2) \triangleq \text{Empty } \Theta_1 \wedge \text{Empty } \Theta_2 \\ & \text{Empty } (\Theta_1 + \Theta_2) \triangleq \text{Empty } \Theta_1 \wedge \text{Empty } \Theta_2 \end{aligned} $
--

■ **Figure 26** Syntactic definition of the *Remove* $\Theta \ h$ operator

1141 The equivalence of the syntactic and semantic definition is captured by Lemma 7.

1142 ► **Lemma 7.** $\llbracket \Theta \rrbracket == \{\}$ if and only if $\text{Empty } \Theta$.

1143 **Proof.** By induction on Θ .

1144 *Case* $\Theta = 0$: We have $\llbracket 0 \rrbracket = \{\}$ and $\text{Empty } 0 = \text{true}$.

1145 *Case* $\Theta = 1$:

1146 We have $\llbracket 1 \rrbracket \neq \{\}$ and $\text{Empty } 1 = \text{false}$.

1147 *Case* $\Theta = h$:

1148 We have $\llbracket h \rrbracket \neq \{\}$ and $\text{Empty } h = \text{false}$.

1149 *Case* $\Theta = \Theta_1 \cdot \Theta_2$: By definition we have $\llbracket \Theta_1 \cdot \Theta_2 \rrbracket = \llbracket \Theta_1 \rrbracket \bullet \llbracket \Theta_2 \rrbracket$ which is equal to
 1150 $\{s_1 \cup s_2 \mid s_1 \in \llbracket \Theta_1 \rrbracket \wedge s_2 \in \llbracket \Theta_2 \rrbracket\}$. It follows that $\llbracket \Theta_1 \cdot \Theta_2 \rrbracket \neq \{\}$ iff $\llbracket \Theta_1 \rrbracket \neq \{\}$ and $\llbracket \Theta_2 \rrbracket \neq \{\}$.
 1151 By induction hypothesis, we have $\llbracket \Theta_1 \rrbracket \neq \{\}$ if and only if $\text{Empty } \Theta_1 = \text{true}$, and $\llbracket \Theta_2 \rrbracket \neq \{\}$ if
 1152 and only if $\text{Empty } \Theta_2 = \text{true}$. The result follows as $\text{Empty } (\Theta_1 \cdot \Theta_2) = \text{Empty } \Theta_1 \wedge \text{Empty } \Theta_2$.

1153 *Case* $\Theta = \Theta_1 + \Theta_2$: By definition we have $\llbracket \Theta_1 + \Theta_2 \rrbracket = \llbracket \Theta_1 \rrbracket \cup \llbracket \Theta_2 \rrbracket$. It follows that
 1154 $\llbracket \Theta_1 + \Theta_2 \rrbracket \neq \{\}$ iff $\llbracket \Theta_1 \rrbracket \neq \{\}$ and $\llbracket \Theta_2 \rrbracket \neq \{\}$. By induction hypothesis, we have $\llbracket \Theta_1 \rrbracket \neq \{\}$ if and
 1155 only if $\text{Empty } \Theta_1 = \text{true}$, and $\llbracket \Theta_2 \rrbracket \neq \{\}$ if and only if $\text{Empty } \Theta_2 = \text{true}$. The result follows
 1156 as $\text{Empty } (\Theta_1 + \Theta_2) = \text{Empty } \Theta_1 \wedge \text{Empty } \Theta_2$. ◀

1157 C Safety of SafeP4

1158 We prove safety in terms of progress and preservation. Both theorems make use of the
 1159 relation $H \models \Theta$ as defined in Figure 17. The empty header instance map only entails the
 1160 empty header instance type 1 (Rule ENT-EMPTY). If a header instance h is contained in
 1161 the map of valid header instances H , H entails the header instance type h (Rule ENT-INST).
 1162 The sequence type $\Theta_1 \cdot \Theta_2$ is entailed by the distinct union of the maps entailing Θ_1 and
 1163 Θ_2 respectively (Rule ENT-SEQ) and the choice type $\Theta_1 + \Theta_2$ is entailed either by the map
 1164 entailing Θ_1 or the map entailing Θ_2 (Rules ENT-CHOICEL and ENT-CHOICER).

1165 We prove progress and preservation only for commands. For expressions we formulate these
 1166 properties as additional lemmas (Lemmas 8 and 9). The respective proofs are straightforward
 1167 for our system.

1168 ► **Lemma 8** (Expression Progress). If $\vdash; \Theta \vdash e : t$ and $H \models \Theta$, then either e is a value or
 1169 $\exists e'. \langle H, e \rangle \rightarrow e'$.

1170 ► **Lemma 9** (Expression Preservation). If $\Gamma; \Theta \vdash e : t$ and $H \models \Theta$ and $\langle H, e \rangle \rightarrow e'$ then
 1171 $\Gamma; \Theta \vdash e' : t$.

1172 ► **Lemma 10** (Expression Substitution). If $\Gamma, x : \tau; \Theta \vdash e : \tau'$ and $\vdash; \cdot \vdash \bar{v} : \bar{\tau}$ then $\Gamma; \Theta \vdash$
 1173 $e[\bar{v}/\bar{x}] : \tau'$

1174 ► **Lemma 11.** If $H \models \Theta$ then $\text{dom}(H) \in \llbracket \Theta \rrbracket$.

1175 **Proof.** By induction on Θ .

1176 *Case* $\Theta = 0$: The case immediately holds as $H \models 0$ is a contradiction.

1177 *Case* $\Theta = 1$: By inversion of *Entailment*, $H = \cdot$, and so $\text{dom}(H) = \{\} \in \llbracket 1 \rrbracket = \{\{\}\}$.

1178 *Case* $\Theta = h$: By inversion of *Entailment*, $\text{dom}(H) = \{h\} \in \llbracket h \rrbracket = \{\{h\}\}$.

1179 *Case* $\Theta = \Theta_1 \cdot \Theta_2$: By inversion of *Entailment*, $H = H_1 \cup H_2$, $H_1 \models \Theta_1$, $H_2 \models \Theta_2$.

1180 By induction hypothesis, $\text{dom}(H_1) \in \llbracket \Theta_1 \rrbracket$ and $\text{dom}(H_2) \in \llbracket \Theta_2 \rrbracket$.

1181 By set theory, $\text{dom}(H) = \text{dom}(H_1) \cup \text{dom}(H_2)$

1182 By induction hypothesis, $\text{dom}(H_1) \in \llbracket \Theta_1 \rrbracket$ and $\text{dom}(H_2) \in \llbracket \Theta_2 \rrbracket$.

1183 By definition of $\llbracket \cdot \rrbracket$ and (\bullet) , $\llbracket \Theta \rrbracket = \llbracket \Theta_1 \rrbracket \bullet \llbracket \Theta_2 \rrbracket = \{hs_1 \cup hs_2 \mid hs_1 \in \llbracket \Theta_1 \rrbracket \wedge hs_2 \in \llbracket \Theta_2 \rrbracket\}$ and

1184 therefore $\text{dom}(H_1) \cup \text{dom}(H_2) \in \{hs_1 \cup hs_2 \mid hs_1 \in \llbracket \Theta_1 \rrbracket \wedge hs_2 \in \llbracket \Theta_2 \rrbracket\}$, i.e., $\text{dom}(H) \in \llbracket \Theta \rrbracket$.

1185 *Case* $\Theta = \Theta_1 + \Theta_2$: By inversion of *Entailment*, either $H \models \Theta_1$ or $H \models \Theta_2$.

1186 *Subcase* $H \models \Theta_1$: By the induction hypothesis, $\text{dom}(H) \in \llbracket \Theta_1 \rrbracket$. and by set theory
1187 $\text{dom}(H) \in \llbracket \Theta_1 \rrbracket \cup \llbracket \Theta_2 \rrbracket$.

1188 *Subcase* $H \models \Theta_2$: Symmetric to the previous subcase. ◀

1189 ▶ **Lemma 12.** *If $H \models \Theta$ and Includes Θ h , then $h \in \text{dom}(H)$.*

1190 **Proof.** By induction on Θ .

1191 *Case* $\Theta = 0$: The case immediately holds as $H \models 0$ is a contradiction.

1192 *Case* $\Theta = 1$: By inversion of *Entailment*, $H = \cdot$. The case immediately holds, as Includes Θ h
1193 is a contradiction.

1194 *Case* $\Theta = g$:

1195 By inversion of *Entailment*, $\text{dom}(H) = \{g\}$

1196 By assumption Includes Θ h , $h = g.\text{Includes } \{g\} g$, i.e., $h \in \text{dom}(H)$.

1197 *Case* $\Theta = \Theta_1 \cdot \Theta_2$:

1198 By inversion of *Entailment*, $H = H_1 \cup H_2$, $H_1 \models \Theta_1$, $H_2 \models \Theta_2$.

1199 By set theory $\text{dom}(H) = \text{dom}(H_1) \cup \text{dom}(H_2)$

1200 By definition of *Inclusion* and by assumption Includes Θ h , Includes Θ_1 $h \vee$ Includes Θ_2 h

1201 *Subcase* Includes Θ_1 h : By induction hypothesis, $h \in \text{dom}(H_1)$ and by assumption
1202 $\text{dom}(H_1) \subseteq \text{dom}(H)$, we can conclude $h \in \text{dom}(H)$.

1203 *Subcase* Includes Θ_2 h : Symmetric to the previous subcase.

1204 *Case* $\Theta = \Theta_1 + \Theta_2$:

1205 By inversion of *Entailment*, either $H \models \Theta_1$ or $H \models \Theta_2$.

1206 By definition of *Inclusion* and by assumption Includes Θ h , Includes Θ_1 h and Includes Θ_2 h .

1207 *Subcase* $H \models \Theta_1$: By induction hypothesis, we can conclude $h \in \text{dom}(H)$.

1208 *Subcase* $H \models \Theta_2$: Symmetric to the previous subcase. ◀

1209 C.1 Control Plane Assumptions

1210 The following propositions model the assumptions about the control plane functions \mathcal{CA} and
1211 \mathcal{CV} that are required to prove type safety.

1212 ▶ **Proposition 13** (Control Plane Reads). If $H \models \Theta$ and $\mathcal{CV}(t) = (\bar{e}, \bar{S})$ and $\Gamma; \Theta \vdash e_i : \tau_i$ for
1213 $e_i \in \bar{e}$ then $\mathcal{CA}(t, H) = (a_i, \bar{v})$.

1214 ▶ **Proposition 14** (Control Plane Action Data). If $H \models \Theta$ and $\mathcal{CA}(t, H) = (a_i, \bar{v})$ and
1215 $\mathcal{A}(a_i) = \lambda \bar{x} : t\bar{a}u. c_i$ then $\cdot; \cdot \vdash \bar{v} : t\bar{a}u$

1216 ▶ **Proposition 15** (Control Plane Assumptions). If $H \models \Theta$ and $\mathcal{CA}(t, H) = (a_i, \bar{v})$ and
1217 $\mathcal{CV}(t) = \bar{S}$ then $H \models \text{Restrict } \Theta \bar{S}$.

1218 C.2 Progress

1219 ▶ **Theorem 16** (Progress). If $\cdot \vdash c : \Theta \Rightarrow \Theta'$ and $H \models \Theta$, then either $c = \text{skip}$ or
1220 $\exists \langle I', O', H', c' \rangle. \langle I, O, H, c \rangle \rightarrow \langle I', O', H', c' \rangle$

1221 **Proof.** By induction on typing derivations of $\cdot \vdash c : \Theta \Rightarrow \Theta'$.

1222 *Case* T-SKIP: $c = \text{skip}$

1223 Immediate.

1224 *Case T-EXTR: $c = \text{extract}(h)$*
 1225 Let $\langle I', v \rangle = \text{deserialize}_\eta(I)$ and $O' = O$ and $H' = H[h \mapsto v]$ and $c' = \text{skip}$. The result
 1226 follows by E-SKIP.
 1227 *Case T-EMIT: $c = \text{emit}(h)$*
 1228 If $h \notin \text{dom}(H)$, let $I' = I$ and $O' = O$ and $H' = H$, and $c' = \text{skip}$. The result follows
 1229 by E-EMITINVALID. Otherwise, $h \in \text{dom}(H)$. Let $H(h) = v$ and $\bar{B} = \text{serialize}_\eta(v)$ and
 1230 $I' = I$ and $O' = O.\bar{B}$ and $H' = H$ and $c' = \text{skip}$. The result follows by E-EMIT.
 1231 *Case T-SEQ: $c = c_1; c_2$ and $\cdot \vdash c_1 : \Theta \Rightarrow \Theta_1$ and $\cdot \vdash c_2 : \Theta_1 \Rightarrow \Theta_2$*
 1232 By induction hypothesis, c_1 is either *skip* or there is some $\langle I', O', H', c'_1 \rangle$, such that
 1233 $\langle I, O, H, c_1 \rangle \rightarrow \langle I', O', H', c'_1 \rangle$.
 1234 If $c = \text{skip}$, let $I' = I$ and $O' = O$ and $H' = H$ and $c' = c_2$. The result follows by E-SEQ.
 1235 Otherwise, the result follows by E-SEQ1.
 1236 *Case T-IF: $c = \text{if } (e) \text{ then } c_1 \text{ else } c_2$ and $\cdot; \Theta \vdash e : \text{Bool}$ and $\cdot \vdash c_1 : \Theta \Rightarrow \Theta_1$ and*
 1237 *$\cdot \vdash c_2 : \Theta \Rightarrow \Theta_2$*
 1238 By the progress theorem for expressions, we have that e is either **true**, **false**, or there is
 1239 some e' such that $\langle H, e \rangle \rightarrow e'$.
 1240 *Subcase $e = \text{true}$:* Let $I' = I$ and $O' = O$ and $H' = H$ and $c' = c_1$. The result follows
 1241 by E-IFTRUE.
 1242 *Subcase $e = \text{false}$:* Symmetric to the previous case.
 1243 *Subcase $\langle H, e \rangle \rightarrow e'$:* Let $I' = I$ and $O' = O$ and $H' = H$ and $c' = \text{if } (e') \text{ } c_1 \text{ } c_2$. The
 1244 result follows by E-IF.
 1245 *Case T-IFVALID: $c = \text{if_valid } (e) \text{ then } c_1 \text{ else } c_2$*
 1246 If $h \in \text{dom}(H)$, let $I' = I$ and $O' = O$ and $H' = H$ and $c' = c_1$. The result follows by
 1247 E-IFVALIDTRUE. Otherwise, $h \notin \text{dom}(H)$. Let $I' = I$ and $O' = O$ and $H' = H$ and
 1248 $c' = c_2$. The result follows by E-IFVALIDFALSE.
 1249 *Case T-APPLY: $c = t.\text{apply}()$*
 1250 By Proposition 13, we have $\mathcal{CA}(t, H) = (a_i, \bar{v})$. Let $\mathcal{A}(a) = \lambda \bar{x} : t.\bar{a}u. c_i$. Let $I' = I$ and
 1251 $O' = O$ and $H' = H$ and $c' = c_i[\bar{v}/\bar{x}]$. The result follows by E-APPLY.
 1252 *Case T-ADD: $c = \text{add}(h)$*
 1253 If $h \in \text{dom}(H)$, let $I' = I$ and $O' = O$ and $H' = H$ and $c' = \text{skip}$. The result follows
 1254 by E-ADDVALID. Otherwise, $h \notin \text{dom}(H)$. Let $v = \text{init}_\eta$ and $I' = I$ and $O' = O$ and
 1255 $H' = H[h \mapsto v]$ and $c' = \text{skip}$. The result follows by E-ADD.
 1256 *Case T-REMOVE: $c = \text{remove}(h)$*
 1257 Let $I' = I$ and $O' = O$ and $H' = H \setminus h$ and $c' = \text{skip}$. The result follows by E-REMOVE.
 1258 *Case T-MOD: $c = h.f = e$ and $\text{Includes } \Theta \text{ } h$ and $\mathcal{F}(h, f) = \tau_i$ and $\cdot; \Theta \vdash e : \tau_i$*
 1259 By the progress rule for expressions, either e is a value or there is some e' such that
 1260 $\langle H, e \rangle \rightarrow e'$.
 1261 *Subcase $e = v$:* By Lemma 12: $h \in \text{dom}(H)$. Let $r = H(h)$ and $r' = \{r \text{ with } f = v\}$.
 1262 Also let $I' = I$ and $O' = O$ and $H' = H[h \mapsto r']$ and $c' = \text{skip}$. The result follows by
 1263 E-MOD.
 1264 *Subcase $\langle H, e \rangle \rightarrow e'$:* Let $I' = I$ and $O' = O$ and $H' = H$ and $c' = h.f = e'$. The result
 1265 follows by E-MOD1.
 1266 *Case T-ZERO: Empty Θ_1*
 1267 By Lemma 11, we have $\text{dom}(H) \in \llbracket \Theta_1 \rrbracket$. By Lemma 7, we have $\llbracket \Theta_1 \rrbracket = \{\}$, which is a
 1268 contradiction. ◀

1269 C.3 Preservation

1270 ► **Lemma 17.** *If $H \models \Theta$ and $h \in \text{dom}(H)$ then $H \models \text{Restrict } \Theta \ h$.*

1271 **Proof.** By induction on Θ .

1272 *Case $\Theta = 0$:* The case immediately holds as $H \models 0$ is a contradiction.

1273 *Case $\Theta = 1$:* By inversion of *Entailment*, $H = \cdot$. The case immediately holds as $h \in \text{dom}(\cdot)$ is a contradiction.

1275 *Case $\Theta = g$:* By inversion of *Entailment*, $\text{dom}(H) = \{g\}$, and so $h = g$.

1276 By definition of *Restriction* $\text{Restrict } \Theta \ h = \text{Restrict } g \ g = g$. By ENT-INST $H \models g$, i.e.,
1277 $H \models \Theta$.

1278 *Case $\Theta = \Theta_1 \cdot \Theta_2$:* By inversion of *Entailment* $H = H_1 \cup H_2, H_1 \models \Theta_1, H_2 \models \Theta_2$. By
1279 $h \in \text{dom}(H)$, either $h \in \text{dom}(H_1)$ or $h \in \text{dom}(H_2)$.

1280 *Subcase $h \in \text{dom}(H_1)$:* By the induction hypothesis, we have $H_1 \models \text{Restrict } \Theta_1 \ h$. By
1281 ENT-SEQ, we have $H_1 \cup H_2 \models \text{Restrict } \Theta_1 \ h \cdot \Theta_2$. By ENT-CHOICE L, we have $H_1 \cup H_2 \models$
1282 $(\text{Restrict } \Theta_1 \ h \cdot \Theta_2) + (\Theta_1 \cdot \text{Restrict } \Theta_2 \ h)$ which finishes the case.

1283 *Subcase $h \in \text{dom}(H_2)$:* Symmetric to the previous subcase.

1284 *Case $\Theta = \Theta_1 + \Theta_2$:* By inversion of *Entailment*, either $H \models \Theta_1$ or $H \models \Theta_2$.

1285 *Subcase $H \models \Theta_1$:* By the induction hypothesis, we have $H \models \text{Restrict } \Theta_1 \ h$. By ENT-
1286 CHOICE L, $H \models \text{Restrict } \Theta_1 \ h + \text{Restrict } \Theta_2 \ h$.

1287 *Subcase $H \models \Theta_2$:* Symmetric to the previous subcase. ◀

1288 ► **Lemma 18** (NegRestrict-Domain-Entail). *If $H \models \Theta$ and $h \notin \text{dom}(H)$ then $H \models$
1289 $\text{NegRestrict } \Theta \ h$.*

1290 **Proof.** By induction on Θ .

1291 *Case $\Theta = 0$:* The case immediately holds as $H \models 0$ is a contradiction.

1292 *Case $\Theta = 1$:* By inversion of *Entailment*, $H = \cdot$. By definition of *Negated Restriction*,
1293 $\text{NegRestrict } \Theta \ h = \text{NegRestrict } 1 \ h = 1$. By ENT-EMPTY $\cdot \models 1$, i.e., $H \models$
1294 $\text{NegRestrict } \Theta \ h$.

1295 *Case $\Theta = g$:* By inversion of *Entailment*, $\text{dom}(H) = \{g\}$. By the induction hypothesis $h \neq g$.
1296 By definition of *Restriction* $\text{NegRestrict } \Theta \ h = \text{NegRestrict } g \ h = g$. By ENT-INST $H \models g$,
1297 i.e., $H \models \text{NegRestrict } \Theta \ h$.

1298 *Case $\Theta = \Theta_1 \cdot \Theta_2$:* By inversion of *Entailment*, $H = H_1 \cup H_2, H_1 \models \Theta_1, H_2 \models \Theta_2$. By
1299 $h \notin \text{dom}(H)$, $h \notin \text{dom}(H_1)$ and $h \notin \text{dom}(H_2)$. By the induction hypothesis, $H_1 \models$
1300 $\text{NegRestrict } \Theta_1 \ h$ and $H_2 \models \text{NegRestrict } \Theta_2 \ h$. By ENT-SEQ, $H_1 \cup H_2 \models \text{NegRestrict } \Theta_1 \ h \cdot$
1301 $\text{NegRestrict } \Theta_2 \ h$ which finishes the case.

1302 *Case $\Theta = \Theta_1 + \Theta_2$:* By inversion of *Entailment*, either $H \models \Theta_1$ or $H \models \Theta_2$.

1303 *Subcase $H \models \Theta_1$:* By the induction hypothesis, we have $H \models \text{NegRestrict } \Theta_1 \ h$. By
1304 Ent-Choice L, $H \models \text{NegRestrict } \Theta_1 \ h + \text{NegRestrict } \Theta_2 \ h$.

1305 *Subcase $H \models \Theta_2$:* Symmetric to the previous subcase. ◀

1306 ► **Lemma 19** (Substitution). *If $\Gamma, x : \tau; \Theta \vdash c : \Theta'$ and $\cdot \vdash v : \tau$ then $\Gamma; \Theta \vdash c : \Theta'$*

1307 **Proof.** By straightforward induction on the derivation $\Gamma, x : \tau; \Theta \vdash c : \Theta'$. ◀

1308 ► **Lemma 20** (Entails-Add). *If $H \models \Theta$ then $H[h \mapsto v] \models \Theta \cdot h$*

1309 **Proof.** We analyze two cases.

1310 *Case $h \in \text{dom}(H)$:* By the assumption of the case, we have $\text{dom}(H) = \text{dom}(H[h \mapsto v])$.

1311 Let $H_1 = H[h \mapsto v]$ and $H_2 = \{h \mapsto v\}$. Observe that $H[h \mapsto v] = H_1 \cup H_2$. By

1312 Lemma 22, we have that $H_1 \models \Theta$. By ENT-INST we have $H_2 \models h$. By ENT-SEQ we have
 1313 $H[h \mapsto v] \models \Theta \cdot h$.

1314 *Case $h \text{ not } \in \text{dom}(H)$:* Let $H_1 = H$ and $H_2 = \{h \mapsto v\}$. Observe that $H[h \mapsto v] = H_1 \cup H_2$.
 1315 By assumption we have $H_1 \models \Theta$. By ENT-INST we have $H_2 \models h$. By ENT-SEQ we have
 1316 $H[h \mapsto v] \models \Theta \cdot h$.

1317 ◀

1318 ► **Lemma 21** (Entails-Remove). *If $H \models \Theta$ then $H \setminus h \models \text{Remove } \Theta \ h$.*

1319 **Proof.** By induction on Θ .

1320 *Case $\Theta = 0$:* The case immediately holds, as $H \models 0$ is a contradiction.

1321 *Case $\Theta = 1$:* By inversion of *Entailment*, $H = \cdot$. By set theory, $\cdot \setminus h = \cdot$ and $\text{Remove } 1 \ h = 1$.
 1322 By ENT-EMPTY, $\cdot \models 1$.

1323 *Case $\Theta = g$:* By inversion of *Entailment*, $\text{dom}(H) = \{g\}$.

1324 *Subcase $g = h$:* By set theory $H \setminus h = \cdot$. By definition of *Remove*, $\text{Remove } \Theta \ h = 1$. By
 1325 ENT-EMPTY, $\cdot \models 1$, which concludes the case.

1326 *Subcase $g \neq h$:* By set theory $H \setminus h = H$. By definition of *Remove*, $\text{Remove } \Theta \ h = g$. By
 1327 assumption, $H \models \Theta$, which concludes the case.

1328 *Case $\Theta = \Theta_1 \cdot \Theta_2$:* By inversion of *Entailment*, $H = H_1 \cup H_2$, $H_1 \models \Theta_1$, $H_2 \models \Theta_2$. By
 1329 induction hypothesis, $H_1 \setminus h \models \text{Remove } \Theta_1 \ h$ and $H_2 \setminus h \models \text{Remove } \Theta_2 \ h$. By set theory,
 1330 $H_1 \setminus h \cup H_2 \setminus h = (H_1 \cup H_2) \setminus h$. By definition of *Removal*, $\text{Remove } \Theta_1 \ h \cdot \text{Remove } \Theta_2 \ h =$
 1331 $\text{Remove } (\Theta_1 \cdot \Theta_2) \ h$. By ENT-SEQ, $(H_1 \cup H_2) \setminus h \models \text{Remove } (\Theta_1 \cdot \Theta_2) \ h$.

1332 *Case $\Theta = \Theta_1 + \Theta_2$:* By inversion of *Entailment*, either $H \models \Theta_1$ or $H \models \Theta_2$.
 1333 By definition of *Removal*, $\text{Remove } \Theta_1 \ h + \text{Remove } \Theta_2 \ h = \text{Remove } (\Theta_1 + \Theta_2) \ h$. *Subcase $H \models \Theta_1$:*
 1334 By induction hypothesis, $H \setminus h \models \Theta_1 \setminus h$. By ENT-CHOICE, applied to $H \setminus h \models \text{Remove } \Theta_1 \ h$,
 1335 and $\text{Remove } \Theta_2 \ h$, we can conclude $H \setminus h \models \text{Remove } \Theta_1 \ h + \text{Remove } \Theta_2 \ h$. By definition of
 1336 *Removal*, $H \setminus h \models \text{Remove } (\Theta_1 + \Theta_2) \ h$. ◀

1337 ► **Lemma 22.** *If $H \models \Theta$ and $\text{dom}(H) = \text{dom}(H')$ then $H' \models \Theta$.*

1338 **Proof.** By induction on Θ .

1339 *Case $\Theta = 0$:* The case immediately holds as $H \models 0$ is a contradiction.

1340 *Case $\Theta = 1$:* By inversion of *Entailment* $H = \cdot$. By assumption $\text{dom}(H) = \text{dom}(H')$, $H' = \cdot$
 1341 and by ENT-EMPTY, $H' \models \Theta$.

1342 *Case $\Theta = g$:* By inversion of *Entailment*, $\text{dom}(H) = \{g\}$. By assumption $\text{dom}(H) = \text{dom}(H')$
 1343 and by Ent-Inst, $H' \models \Theta$.

1344 *Case $\Theta = \Theta_1 \cdot \Theta_2$:* By inversion of *Entailment*, $H = H_1 \cup H_2$, $H_1 \models \Theta_1$, $H_2 \models \Theta_2$. By set
 1345 theory, $\text{dom}(H) = \text{dom}(H_1) \cup \text{dom}(H_2)$. By induction hypothesis if $\text{dom}(H'_1) = \text{dom}(H_1)$
 1346 and $\text{dom}(H'_2) = \text{dom}(H_2)$, then $H'_1 \models \Theta_1$, $H'_2 \models \Theta_2$. By ENT-SEQ, $H' = H'_1 \cup H'_2 \models \Theta_1 \cdot \Theta_2$.

1347 *Case $\Theta = \Theta_1 + \Theta_2$:* By inversion of *Entailment*, either $H \models \Theta_1$ or $H \models \Theta_2$.

1348 *Subcase $H \models \Theta_1$:* By induction hypothesis, we have $H' \models \Theta_1$. By ENT-CHOICE, $H' \models$
 1349 $\Theta_1 + \Theta_2$.

1350 *Subcase $H \models \Theta_2$:* Symmetric to the previous subcase. ◀

1351 We define $\Theta_1 < \Theta_2 \triangleq \llbracket \Theta_1 \rrbracket \subseteq \llbracket \Theta_2 \rrbracket$, i.e., for every $S \in \llbracket \Theta_1 \rrbracket$, $S \in \llbracket \Theta_2 \rrbracket$.

1352 ► **Lemma 23.** *If $\Theta'_1 < \Theta_1$ then $\Theta_1 \cdot h < \Theta_1 \cdot h$.*

1353 **Proof.** We calculate as follows:

1354 1. $\llbracket \Theta'_1 \rrbracket \subseteq \llbracket \Theta_1 \rrbracket$ by (B) and the definition of $<$.

- 1355 2. $\llbracket \Theta'_1 \cdot h \rrbracket == \llbracket \Theta'_1 \rrbracket \bullet \{\{h\}\}$ by definition of $\llbracket \cdot \rrbracket$
 1356 3. $\llbracket \Theta_1 \cdot h \rrbracket == \llbracket \Theta_1 \rrbracket \bullet \{\{h\}\}$ by definition of $\llbracket \cdot \rrbracket$
 1357 4. Let $S \in \llbracket \Theta'_1 \rrbracket \bullet \{\{h\}\}$.
 1358 5. $S = S' \cup \{h\}$, where $S' \in \llbracket \Theta'_1 \rrbracket$ by def of \bullet .
 1359 6. By 1., $S' \in \llbracket \Theta_1 \rrbracket$
 1360 7. By set theory, $S' \cup \{h\} \in \llbracket \Theta_1 \rrbracket \bullet \{\{h\}\}$.
 1361 8. Then $\llbracket \Theta'_1 \rrbracket \bullet \{\{h\}\} \subseteq \llbracket \Theta_1 \rrbracket \bullet \{\{h\}\}$

1362

1363 ► **Lemma 24.** *If $\Theta'_1 < \Theta_1$ then $\llbracket \text{Remove } \Theta'_1 h \rrbracket \subseteq \llbracket \text{Remove } \Theta_1 h \rrbracket$.*

1364 **Proof.** Since $\llbracket \text{Remove } \Theta'_1 h \rrbracket == \llbracket \Theta'_1 \rrbracket \setminus h$ and $\llbracket \text{Remove } \Theta_1 h \rrbracket == \llbracket \Theta_1 \rrbracket \setminus h$ by Lemma 6, we
 1365 can equivalently show that $\llbracket \Theta'_1 \rrbracket \setminus h \subseteq \llbracket \Theta_1 \rrbracket \setminus h$, which follows from set theory. ◀

1366 ► **Lemma 25.** *If $\Theta'_1 < \Theta_1$ then $\llbracket \text{Restrict } \Theta'_1 h \rrbracket \subseteq \llbracket \text{Restrict } \Theta_1 h \rrbracket$*

1367 **Proof.** By Lemma 3, $\llbracket \text{Restrict } \Theta'_1 h \rrbracket == \llbracket \Theta'_1 \rrbracket | h$ and $\llbracket \text{Restrict } \Theta_1 h \rrbracket == \llbracket \Theta_1 \rrbracket | h$. By set
 1368 theory, $\llbracket \Theta'_1 \rrbracket | h \subseteq \llbracket \Theta_1 \rrbracket | h$ when $\llbracket \Theta'_1 \rrbracket \subseteq \llbracket \Theta_1 \rrbracket$, so we are done. ◀

1369 ► **Lemma 26.** *If $\Theta'_1 < \Theta_1$ then $\llbracket \text{NegRestrict } \Theta'_1 h \rrbracket \subseteq \llbracket \text{NegRestrict } \Theta_1 h \rrbracket$*

1370 **Proof.** By Lemma 4, $\llbracket \text{NegRestrict } \Theta'_1 h \rrbracket == \llbracket \Theta'_1 \rrbracket | \neg h$ and $\llbracket \text{NegRestrict } \Theta_1 h \rrbracket ==$
 1371 $\llbracket \Theta_1 \rrbracket | \neg h$. By set theory, $\llbracket \Theta'_1 \rrbracket | \neg h \subseteq \llbracket \Theta_1 \rrbracket | \neg h$ when $\llbracket \Theta'_1 \rrbracket \subseteq \llbracket \Theta_1 \rrbracket$, so we are done. ◀

1372 ► **Lemma 27.** *If $\Theta' < \Theta$ and $\text{Includes } \Theta_1 h$ then $\text{Includes } \Theta'_1 h$.*

1373 **Proof.** By Lemma 5, $\text{Includes } \Theta'_1 h = h \sqsubset \Theta'_1$. By the same lemma, $\text{Includes } \Theta_1 h = h \sqsubset$
 1374 Θ_1 . Let $S \in \llbracket \Theta'_1 \rrbracket$ to show $h \in S$ and hence $h \sqsubset \Theta'_1$. Since $\llbracket \Theta'_1 \rrbracket \subseteq \llbracket \Theta_1 \rrbracket$, by assumption and
 1375 definition of $<$, then $S \in \llbracket \Theta_1 \rrbracket$. Since $h \sqsubset \Theta_1$, conclude $h \in S$ and we are done. ◀

1376 ► **Lemma 28.** *If $\Theta'_a < \Theta_a$ and $\Theta'_b < \Theta_b$ then $\Theta'_a + \Theta'_b < \Theta_a + \Theta_b$.*

1377 **Proof.** We have to show that $\llbracket \Theta'_a + \Theta'_b \rrbracket \subseteq \llbracket \Theta_a + \Theta_b \rrbracket$ when $\Theta'_a < \Theta_a$ and $\Theta'_b < \Theta_b$. By
 1378 definition of $\llbracket \cdot \rrbracket$ we can equally show that $\llbracket \Theta'_a \rrbracket \cup \llbracket \Theta'_b \rrbracket \subseteq \llbracket \Theta_a \rrbracket \cup \llbracket \Theta_b \rrbracket$, which follows from set
 1379 theory. ◀

1380 ► **Lemma 29.** *If $\Gamma; \Theta_1 \vdash a : \bar{\tau} \rightarrow \Theta_2$ and $\Theta'_1 < \Theta_1$, then $\exists \Theta'_2. \Gamma; \Theta'_1 \vdash a : \bar{\tau} \rightarrow \Theta'_2$ and*
 1381 $\Theta'_2 < \Theta_2$.

1382 **Proof.** There is only one way to have concluded that $\Gamma; \Theta_1 \vdash a : \bar{\tau} \rightarrow \Theta_2$: via the [T-ACTION]
 1383 rule, which gives us two facts: we know $a - \lambda \bar{x} : \bar{\tau}. c$, and $\Gamma, \bar{x} : \bar{\tau} \vdash \Theta_2 : \Theta_1 \Rightarrow$.

1384 Since this c is an action command, is only generated by the add, remove, modification
 1385 and sequence commands. So we perform a limited induction on the structure of c :

1386 *Case $c = \text{add}(h)$.* The only typing rule that applies is T-ADD, so we know $\Theta_2 = \Theta_1 \cdot h$.
 1387 Now let $\Theta'_2 = \Theta'_1 \cdot h$. Then T-ADD shows $\Gamma, \bar{x} : \bar{\tau} \vdash \text{add}(h) : \Theta'_1 \Rightarrow \Theta'_1 \cdot h$. Then $\Theta'_1 \cdot h < \Theta_1 \cdot h$
 1388 follows by Lemma 23, and we are done.

1389 *Case $c = \text{remove}(h)$.* The only typing rule that could have applied is T-REMOVE,
 1390 so we know that $\Theta_2 = \text{Remove } \Theta_1 h$. Let $\Theta'_2 = \text{Remove } \Theta'_1 h$. Then T-REMOVE shows
 1391 $\Gamma, \bar{x} : \bar{\tau} \vdash \text{remove}(h) : \Theta'_1 \Rightarrow \text{Remove } \Theta'_1 h$. Then $\text{Remove } \Theta'_1 h < \text{Remove } \Theta_1 h$ by Lemma 24.

1392 *Case $c = h.f = v$.* The only typing rule that could have applied is T-MOD, so we know
 1393 that $\Theta_2 = \Theta_1$, Let $\Theta'_2 = \Theta'_1$, which proves $\Theta'_2 < \Theta_2$ by assumption.

1394 We know by our case assumption that $\Gamma, \bar{x} : \bar{\tau}; \Theta_1 \vdash e : \mathcal{F}(h, f)$ and **Includes** $\Theta_1 h$. By
 1395 T-MOD, we only need to show that (1) $\Gamma, \bar{x} : \bar{\tau} \vdash e : \mathcal{F}(h, f)$ and (2) **Includes** $\Theta'_1 h$. (1)
 1396 follows by Lemma 30, and (2) follows by Lemma 27.

1397 *Case* $c = c_1; c_2$. The only rule that could have applied is T-SEQ, so we know that
 1398 $\Gamma, \bar{x} : \bar{\tau} \vdash c_1 : \Theta_1 \Rightarrow \Theta_{11}$, and $\Gamma, \bar{x} : \bar{\tau} \vdash c_2 : \Theta_{11} \Rightarrow \Theta_2$, and $\Theta'_1 \leq \Theta_1$.

1399 The inductive hypothesis on c_1 gives us a $\Theta'_{11} < \Theta_{11}$ such that $\Gamma, \bar{x} : \bar{\tau} \vdash c_1 : \Theta'_{11} \Rightarrow T'_{11}$.

1400 The inductive hypothesis on c_2 gives us a $\Theta'_2 < \Theta_2$ such that $\Gamma, \bar{x} : \bar{\tau} \vdash c_2 : \Theta'_{11} \Rightarrow \Theta'_2$.

1401 The result follows by T-Seq.

1402

1403 ► **Lemma 30.** *If $\Gamma; \Theta \vdash e : \tau$ and $\Theta' < \Theta$, then $\Gamma; \Theta' \vdash e : \tau$.*

1404 **Proof.** By induction on the typing derivation.

1405 *Case* T-CONSTANT We know $e = k(\bar{e})$, $\Gamma; \Theta \vdash e_i : \tau_i$ for all i , $\text{typeof}(k) = t\bar{a}u \rightarrow \tau$ and
 1406 $\Theta' < \Theta$. By induction hypothesis, $\Gamma; \Theta' \vdash e_i : \tau_i$ for all i and we are done by T-CONSTANT.

1407 *Case* T-VAR We know $e = x$, $x : \tau \in \Gamma$, and $\Theta' < \Theta$. We are done by T-VAR.

1408 *Case* T-FIELD We know $e = h.f$, **Includes** Θh and $\Theta' < \Theta$. By Lemma 27, we know
 1409 **Includes** $\Theta' h$ and the result follows by T-FIELD. ◀

1410 ► **Lemma 31.** *If $\Gamma \vdash c : \Theta_1 \Rightarrow \Theta_2$ and $\Theta'_1 < \Theta_1$, then $\exists \Theta'_2. \Gamma \vdash c : \Theta'_1 \Rightarrow \Theta'_2$ and $\Theta'_2 < \Theta_2$.*

1411 **Proof.** By induction on a derivation of $\Gamma \vdash c : \Theta_1 \Rightarrow \Theta_2$. We refer to assumptions $\Gamma \vdash c :$
 1412 $\Theta_1 \Rightarrow \Theta_2$ and $\Theta'_1 < \Theta_1$ as (A) and (B) respectively. Similarly, we use (1) and (2) to refer to
 1413 the proof goals $\exists \Theta'_2. \Gamma \vdash c : \Theta'_1 \Rightarrow \Theta'_2$ and $\Theta'_2 < \Theta_2$ respectively.

1414 *Case* T-ZERO: By assumption, we have **Empty** Θ_1 . By Lemmas 7 and 27 we have
 1415 **Empty** Θ'_1 . Let $\Theta'_2 = \Theta_2$. We have $\Gamma \vdash c : \Theta'_1 \Rightarrow \Theta_2$ by T-ZERO, proving (1), and $\Theta'_2 < \Theta_2$ by
 1416 reflexivity, proving (2).

1417 *Case* T-SKIP: We know $c = \text{skip}$ and $\Theta_2 = \Theta_1$ and $\Theta'_1 < \Theta_1$. Let $\Theta'_2 = \Theta'_1$. Then by
 1418 assumption (B) $\Theta'_2 = \Theta'_1 < \Theta_1 = \Theta_2$, proving (2) and $\Gamma \vdash \text{skip} : \Theta'_1 \Rightarrow \Theta'_1$ by T-SKIP, proving
 1419 (1).

1420 *Case* T-EMIT: We know $c = \text{emit}(h)$ and $\Theta_2 = \Theta_1$ and $\Theta'_1 < \Theta_1$. Let $\Theta'_2 = \Theta'_1$. Then by
 1421 assumption (B), $\Theta'_2 = \Theta'_1 < \Theta_1 = \Theta_2$, proving (2) and $\Gamma \vdash \text{emit}(h) : \Theta'_1 \Rightarrow \Theta'_1$ by T-EMIT,
 1422 proving (1).

1423 *Case* T-ADD: We know $c = \text{add}(h)$ and $\Theta_2 = \Theta_1 \cdot h$ and $\Theta'_1 < \Theta_1$. (1) follows since we
 1424 can prove $\Gamma \vdash \text{add}(h) : \Theta'_1 \Rightarrow \Theta'_1 \cdot h$ by T-ADD. (2), i.e., $\Theta'_1 \cdot h < \Theta_1 \cdot h$, follows from Lemma
 1425 23.

1426 *Case* T-EXTR: Similar to case T-ADD. We know $c = \text{extract}(h)$ and $\Theta_2 = \Theta_1 \cdot h$ and
 1427 $\Theta'_1 < \Theta_1$. Let $\Theta'_2 = \Theta'_1 \cdot h$. (1) follows since we can prove $\Gamma \vdash \text{extract}(h) : \Theta'_1 \Rightarrow \Theta'_1 \cdot h$ by
 1428 T-EXTRACT. (2) follows by Lemma 23.

1429 *Case* T-REM: We know $c = \text{remove}(h)$ and $\Theta_2 = \text{Remove } \Theta_1 h$ and $\Theta'_1 < \Theta_1$. Let
 1430 $\Theta'_2 = \text{Remove } \Theta'_1 h$. (1) follows by T-REM and for (2) we have to show that $\text{Remove } \Theta'_1 h <$
 1431 $\text{Remove } \Theta_1 h$, which follows from Lemma 24.

1432 *Case* T-MOD: We know $c = h.f = e$ and $\Theta_2 = \Theta_1$ and $\Theta'_1 < \Theta_1$. Let $\Theta'_2 = \Theta'_1$. If
 1433 $\llbracket \Theta'_1 \rrbracket = \llbracket 0 \rrbracket$ then $\Theta_1 == 0$ by idempotent semiring equality and (1) follows by T-ZERO.
 1434 Otherwise $\llbracket \Theta'_1 \rrbracket$ is nonempty. To show (1) we need to show

1435 (a) **Includes** $\Theta'_1 h$,

1436 (b) $\mathcal{F}(h, f) = \tau$,

1437 (c) $\Gamma; \Theta_1 \vdash e : \tau$

(b) and (c) follow from the assumption that the previous rule in the typing derivation was T-MOD. This inversion also gives us **Includes** $\Theta_1 h$. To show (a) we calculate as follows. $h \sqsubset \llbracket \Theta_1 \rrbracket$ by Lemma 5, i.e. $h \in S$ for every $S \in \llbracket \Theta_1 \rrbracket$ by definition. Since $\llbracket \Theta'_1 \rrbracket \subseteq \llbracket \Theta_1 \rrbracket$, $h \in S$ for every $S \in \llbracket \Theta'_1 \rrbracket$, by set theory. By definition we get $h \sqsubset \llbracket \Theta'_1 \rrbracket$. By Lemma 5, we can conclude **Includes** $\Theta'_1 h$.

Case T-SEQ: We know $c = c_1; c_2$ and $\Gamma \vdash c_1 : \Theta_1 \Rightarrow \Theta_{11}$ and $\Gamma \vdash c_2 : \Theta_{11} \Rightarrow \Theta_2$ and $\Theta'_1 < \Theta_1$. By induction hypothesis, $\exists \Theta'_{11}. \Gamma \vdash c_1 : \Theta'_1 \Rightarrow \Theta'_{11}$ and $\Theta'_{11} < \Theta_{11}$. Again, by induction hypothesis, $\exists \Theta'_2. \Gamma \vdash c_2 : \Theta'_{11} \Rightarrow \Theta'_2$ (proving 1) and $\Theta'_2 < \Theta_2$ (proving 2) which concludes the case.

Case T-IFVALID: We know $c = \text{valid}(h) \ c_1 \ \text{else} \ c_2$ and $\Gamma \vdash c_1 : \text{Restrict } \Theta_1 h \Rightarrow \Theta_t, \Gamma \vdash c_2 : \text{NegRestrict } \Theta_1 h \Rightarrow \Theta_f, \Theta_2 = \Theta_t + \Theta_f$, and $\Theta'_1 < \Theta_1$. Let $\Theta'_2 = \text{Restrict } \Theta'_1 h + \text{NegRestrict } \Theta'_1 h$. (1) is immediate from T-IFVALID. (2) follows from Lemmas 25, 26 and 28.

Case T-IF: We know $c = \text{if } (e) \ c_1 \ \text{else} \ c_2, \Gamma \vdash c_1 : \Theta_1 \Rightarrow \Theta_{11}, \Gamma \vdash c_2 : \Theta_1 \Rightarrow \Theta_{12}, \Gamma; \Theta_1 \vdash e : \text{Bool}$ and $\Theta'_1 < \Theta_1$. By induction hypothesis, there exists Θ'_{11} such that (1a) $\Gamma \vdash c_1 : \Theta'_1 \Rightarrow \Theta'_{11}$ and (2a) $\Theta'_{11} < \Theta_{11}$. Also by induction hypothesis, there exists Θ'_{12} such that (1b) $\Gamma \vdash c_2 : \Theta'_1 \Rightarrow \Theta'_{12}$ and (2b) $\Theta'_{12} < \Theta_{12}$. Let $\Theta'_2 = \Theta'_{11} + \Theta'_{12}$. (1) follows by T-IF (1a), (1b), and the fact that $\Gamma; \Theta_1 \vdash e : \text{Bool}$. (2) follows by Lemma 28, (2a), (2b).

Case T-APPLY: We know $c = t.\text{apply}(), \Theta_2 = \Theta_{11} + \Theta_{12} + \dots + \Theta_{1n}, t.\text{actions} = a_1 + a_2 + \dots + a_n, \vdash; \Theta_1 \vdash e_j : \tau_j$ for $j = 1, \dots, m, \mathcal{V}(t) = (S_1 \dots S_n, e_1 \dots e_m)$ and $\text{Restrict } \Theta_1 S_i \vdash a_i : \bar{\tau}_i \rightarrow \Theta_{1i}$. We want to construct $\Theta'_2 < \Theta_2$ such that $\Gamma \vdash t.\text{apply}() : \Theta'_1 \Rightarrow \Theta'_2$. By repeated application of Lemma 25, $\text{Restrict } \Theta'_1 S_i < \text{Restrict } \Theta_1 S_i$. For every i apply Lemma 29 which gives us $\Gamma; \text{Restrict } \Theta'_1 S_i \vdash a : \bar{\tau} \rightarrow \Theta_{1i}$ and $\Theta'_{1i} < \Theta_{1i}$. Let $\Theta'_2 = \sum_i \Theta'_{1i}$. (2) follows by T-APPLY. To show (1), i.e., $\Theta'_2 = \sum_i \Theta'_{1i} < \sum_i \Theta_{1i} = \Theta_2$. We know $\Theta'_{1i} < \Theta_{1i}$ for all i . The result follows by repeated application of Lemma 28. \blacktriangleleft

► **Theorem 32 (Preservation).** *If $\Gamma \vdash c : \Theta_1 \Rightarrow \Theta_2$ and $\langle I, O, H, c \rangle \rightarrow \langle I', O', H', c' \rangle$, where $H \models \Theta_1$, then $\exists \Theta'_1, \Theta'_2. \Gamma \vdash c : \Theta'_1 \Rightarrow \Theta'_2$ where $H' \models \Theta'_1$ and $\Theta'_2 < \Theta_2$.*

Proof. By induction on a derivation of $\Gamma \vdash c : \Theta_1 \Rightarrow \Theta_2$, with a case analysis on the last rule used.

Case T-SKIP: $c = \text{skip}$ and $\Theta_2 = \Theta_1$

Vauously holds as there is no c' such that $\langle I, O, H, c \rangle \rightarrow \langle I', O', H', c' \rangle$.

Case T-EXTR: $c = \text{extract}(h)$ and $\Theta_2 = \Theta_1 \cdot h$

The only evaluation rule that applies to c is E-EXTR, so we also have $c' = \text{skip}$ and $\mathcal{HT}(h) = \eta$ and $H' = H[h \mapsto v]$ where $\text{deserialize}_\eta = (v, I')$. Let $\Theta'_1 = \Theta'_2 = \Theta_2$. We have $\Gamma \vdash c' : \Theta'_1 \Rightarrow \Theta'_2$ by T-SKIP, we have $H' \models \Theta'_2$ by Lemma 20, and we have $\Theta'_2 < \Theta_2$ by reflexivity.

Case T-EMIT: $c = \text{emit}(h)$ and $\Theta_2 = \Theta_1$.

There are two evaluation rules that apply to c , E-EMIT and E-EMITINVALID. In either case, $c' = \text{skip}$ and $H' = H$. Let $\Theta'_1 = \Theta'_2 = \Theta_1$. We have $\Gamma \vdash c' : \Theta'_1 \Rightarrow \Theta'_2$ by T-SKIP, we have $H' \models \Theta'_1$ by assumption, and we have $\Theta'_2 < \Theta_2$ by reflexivity.

Case T-SEQ: $c = c_1; c_2$ and $\Gamma \vdash c_1 : \Theta_1 \Rightarrow \Theta_{12}$ and $\Gamma \vdash c_2 : \Theta_{12} \Rightarrow \Theta_2$

There are two evaluation rules that apply to c , E-SEQ1 and E-SEQ.

Subcase E-SEQ: $c' = c_2$ and $H' = H$

By inversion of $\Gamma \vdash c_1 : \Theta_1 \Rightarrow \Theta_{12}$ we have $\Theta_{12} = \Theta_1$. Let $\Theta'_1 = \Theta_1$ and $\Theta'_2 = \Theta_2$. We have $\Gamma \vdash c' : \Theta'_1 \Rightarrow \Theta'_2$ by assumption, we have $H \models \Theta'_1$ also by assumption, and $\Theta'_2 < \Theta_2$ by reflexivity.

1484 **Subcase E-SEQ:** $c' = c'_1; c_2$ and $\langle I, O, H, c_1 \rangle \rightarrow \langle I', O', H', c'_1 \rangle$.

1485 By IH we have $\Gamma \vdash c_1 : \Theta'_1 \Rightarrow \Theta'_{12}$ such that $H' \models \Theta'_1$ and $\Theta'_{12} < \Theta_{12}$. By Lemma 31

1486 we have $\Gamma \vdash c_2 : \Theta'_{12} \Rightarrow \Theta_2$ for some $\Theta'_2 < \Theta_2$. We have $\Gamma \vdash c_1; c_2 : \Theta'_1 \Rightarrow \Theta'_2$ by T-SEQ,

1487 which finishes the case.

1488 **Case T-IF:** $c = \text{if } (e) \ c_1 \ \text{else } c_2$ and $\Gamma; \Theta_1 \vdash e : \text{Bool}$ and $\Gamma \vdash c_1 : \Theta_1 \Rightarrow \Theta_{12}$ and

1489 $\Gamma \vdash c_2 : \Theta_1 \Rightarrow \Theta_{22}$ and $\Theta_2 = \Theta_{12} + \Theta_{22}$.

1490 There are two evaluation rules that apply to c , E-IF, E-IFTRUE, and E-IFFALSE.

1491 **Subcase E-IF:** $c' = \text{if } (e') \ c_1 \ \text{else } c_2$ and $H' = H$

1492 Let $\Theta'_1 = \Theta_1$ and $\Theta'_2 = \Theta_2$. We have $\Gamma \vdash \text{if } e \ c_1 \ c_2 : \Theta'_1 \Rightarrow \Theta'_2$ by T-IF, we have

1493 $H \models \Theta_1$ by assumption, and we have $\Theta_2 < \Theta'_2$ by reflexivity.

1494 **Subcase E-IFTRUE:** $c' = c_1$ and $H' = H$.

1495 Let $\Theta'_1 = \Theta_1$ and $\Theta'_2 = \Theta_{12}$. We have $\Gamma \vdash c' : \Theta'_1 \Rightarrow \Theta'_2$ by assumption, we have

1496 $H \models \Theta'_1$ also by assumption, and we have $\Theta'_2 < \Theta_2$ by the definition of $<$ and the

1497 semantics of types.

1498 **Subcase E-IFTRUE:** $c' = c_2$ and $H' = H$.

1499 Symmetric to the previous case.

1500 **Case T-IFVALID:** $c = \text{valid}(h) \ c_1 \ \text{else } c_2$ and $\Gamma \vdash c_1 : \text{Restrict } \Theta_1 \ h \Rightarrow \Theta_{12}$ and $\Gamma \vdash c_2 :$

1501 $\text{NegRestrict } \Theta_1 \ h \Rightarrow \Theta_{22}$ and $\Theta_2 = \Theta_{12} + \Theta_{22}$.

1502 There are two evaluation rules that apply to c , E-IFVALIDTRUE and E-IFVALIDFALSE

1503 **Subcase E-IFVALIDTRUE:** $c' = c_1$ and $h \in \text{dom}(H)$ and $H' = H$.

1504 Let $\Theta'_1 = \text{Restrict } \Theta_1 \ h$ and $\Theta'_2 = \Theta_{12}$. We have $\Gamma \vdash c' : \Theta'_1 \Rightarrow \Theta'_2$ by assumption,

1505 we have $H \models \Theta'_1$ by Lemma 17, and we have $\Theta'_2 < \Theta_2$ by the definition of $<$ and

1506 semantics of types.

1507 **Subcase E-IFVALIDFALSE:** $c' = c_2$ and $h \notin \text{dom}(H)$ and $H' = H$.

1508 Symmetric to the previous case.

1509 **Case T-APPLY:** $c = t.\text{apply}()$ and $\mathcal{CV}(t) = (\bar{S}, \bar{e})$ and $t.\text{actions} = \bar{a}$ and $;\Theta \vdash e_i : \tau_i$ for

1510 $e_i \in \bar{e}$ and $\text{Restrict } \Theta_1 \ S_i \vdash a_i : \bar{\tau}_i \rightarrow \Theta'_i$ for $a_i \in \bar{a}$ and $\Theta_2 = \sum (\Theta'_i)$

1511 There is only one evaluation rule that applies to c , E-APPLY. It follows that $\mathcal{CA}(H, t) =$

1512 (a_i, \bar{v}) , and $c' = c_i[\bar{v}/\bar{x}]$ where $\mathcal{A}(a_i) = \lambda \bar{x}. c_i$. Next, inverting T-Action, we have

1513 $\Gamma, \bar{x} : \bar{\tau}_i; \vdash c_i : \text{Restrict } \Theta \ S_i \Rightarrow \Theta'_i$. By Proposition 14, we have $;\cdot \vdash \bar{v} : \bar{\tau}_i$. Hence, by the

1514 substitution lemma, we have $\Gamma \vdash c_i[\bar{v}/\bar{x}] : \text{Restrict } \Theta \ S_i \Rightarrow \Theta'_i$. Let $\Theta'_1 = \text{Restrict } \Theta \ S_i$

1515 and $\Theta'_2 = \Theta'_i$. We have already shown that $\Gamma \vdash c' : \Theta'_1 \Rightarrow \Theta'_2$, we have that $H' \models \Theta'_1$ by

1516 Proposition 15, and we have $\Theta'_2 < \Theta_2$ by the definition of $<$ and the semantics of union

1517 types.

1518 **Case T-ADD:** $c = \text{add}(h)$ and $\Theta_2 = \Theta_1 \cdot h$

1519 There are two evaluation rules that apply to c , E-ADD and E-ADDVALID.

1520 **Subcase E-ADD:** $c' = \text{skip}$ and $\mathcal{HT}(h) = \eta$ and $\text{init}_\eta = v$ and $H' = H[h \mapsto v]$

1521 Let $\Theta'_1 = \Theta'_2 = \Theta_2$. We have $\Gamma \vdash c' : \Theta'_1 \Rightarrow \Theta'_2$ by T-SKIP, we have $H' \models \Theta'_1$ by

1522 Lemma 20, and we have $\Theta'_2 < \Theta_2$ by reflexivity.

1523 **Subcase E-ADDVALID:** $c' = \text{skip}$ and $H' = H$

1524 Let $\Theta'_1 = \Theta'_2 = \Theta_2$. We have $\Gamma \vdash c' : \Theta'_1 \Rightarrow \Theta'_2$ by T-SKIP, we have $H' \models \Theta'_1$ by

1525 Lemma 22 and 20 since $\text{dom}(H') = \text{dom}(H[h \mapsto v])$ for any v , and we have $\Theta'_2 < \Theta_2$

1526 by reflexivity.

1527 *Case T-REM: $c = \text{remove}(h)$ and $\Theta_2 = \text{Remove } \Theta_1 \ h$*
 1528 There is only one evaluation rule that applies to c , E-REM, so we have $c' = \text{skip}$ and
 1529 $H' = H \setminus h$. Let $\Theta'_1 = \Theta'_2 = \text{Remove } \Theta \ h$. We have $\Gamma \vdash c' : \Theta'_1 \Rightarrow \Theta'_2$ by T-SKIP, we have
 1530 $H' \models \Theta'_1$ by Lemma 21, and we have $\Theta'_2 < \Theta_2$ by reflexivity.
 1531 *Case T-MOD: $c = h.f = e$ and $\text{Includes } \Theta_1 \ h$ and $\mathcal{HT}(h, f) = \tau_i$ and $\cdot; \Theta_1 \vdash e : \tau_i$ and*
 1532 $\Theta_2 = \Theta_1$
 1533 There are two evaluation rules that applies to c , E-MOD1 and E-MOD.
 1534 **Subcase E-MOD1: $c' = h.f = e'$ and $e \rightarrow e'$ and $H' = H$**
 1535 By preservation for expressions we have $\cdot; \Theta_1 \vdash e' : \tau_i$. Let $\Theta'_1 = \Theta'_2 = \Theta_1$. We have
 1536 $\Gamma \vdash c' : \Theta'_1 \Rightarrow \Theta'_2$ by T-MOD, we have $H' \models \Theta'_1$ by assumption, and we have $\Theta'_2 < \Theta_2$
 1537 by reflexivity.
 1538 **Subcase E-MOD: $c' = \text{skip}$ and $\text{dom}(H') = \text{dom}(H)$**
 1539 Let $\Theta'_1 = \Theta'_2 = \Theta_1$. We have $\Gamma \vdash c' : \Theta'_1 \Rightarrow \Theta'_2$ by T-SKIP, we have $H' \models \Theta'_1$ by
 1540 Lemma 22, and we have $\Theta'_2 < \Theta_2$ by reflexivity.
 1541 *Case T-ZERO: Empty Θ_1*
 1542 By Lemma 11, we have $\text{dom}(H) \in \llbracket \Theta_1 \rrbracket$. By Lemma 7, we have $\llbracket \Theta_1 \rrbracket = \{\}$, which is a
 1543 contradiction. ◀